



LEADS 3-Day Camp

Session 3: Data Management:

Relational Databases and SQL

Il-Yeol Song, Ph.D.

Professor

College of Computing & Informatics

Drexel University

Philadelphia, PA 19104

Song@drexel.edu

<http://www.cci.drexel.edu/faculty/song/>

Il-Yeol Song, Ph.D.

1

1

Hello!



2

1

Instructor

- Professor, College of Computing & Informatics, Drexel University,
- PhD in CS, LSU, Baton Rouge, USA, 1988
- PhD Program Director (2010-2015)
- CVDI Deputy Director (2012-2014)
- Research Topics: *Conceptual Modeling, Data Warehousing, Big Data Management and Data Analytics, Smart Aging*
- Elected as an **ER Fellow**, 2012
- Named an **ACM Distinguished Scientist** in 2013
- Received **Peter Chen Award in Conceptual Modeling** in 2015
- Four teaching awards from Drexel (1991, 2000, 2001, 2011) including **Lindback Distinguished Teaching Award** (2001)
- Co-Editor-in-Chief, Journal of Computing Science & Engineering
- Consulting Editor, Data & Knowledge Engineering Journal
- Co-Chair, iSchool Model Data Science Curriculum Committee
- Chair, IEEE Big Data and Smart Computing (BigComp) Conference
- Published about 200+ papers



Il-Yeol Song, Ph.D.

3

Acknowledgement

- Materials from:
 - Drexel CCI INFO 605 and 606 Class (Il-Yeol Song)
 - Some diagrams and examples from
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management, 13th Edition, 2017. Cengage Learning.



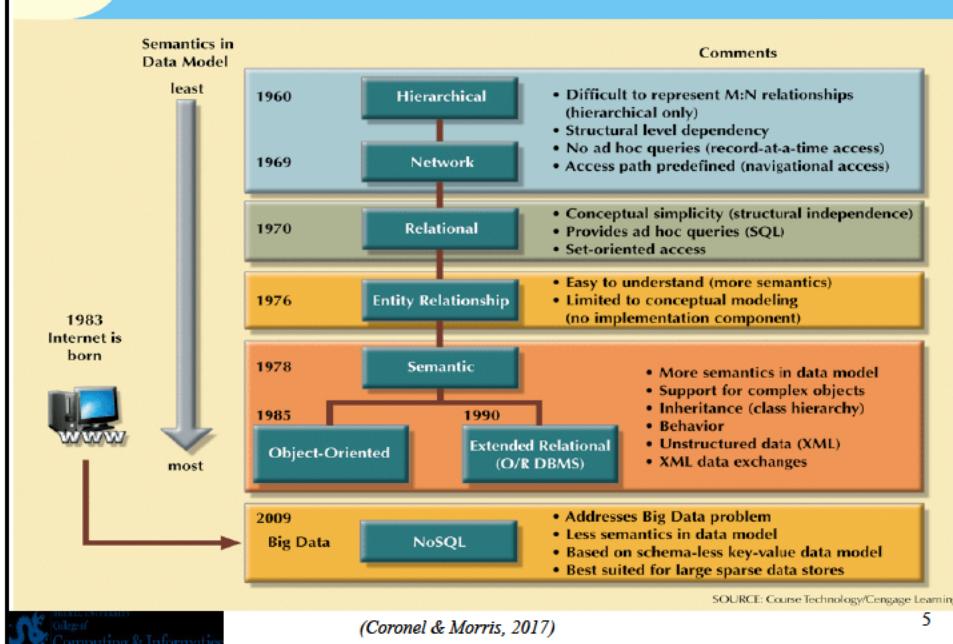
Il-Yeol Song, Ph.D.

4

4

FIGURE
2.6

The evolution of data models



5

5

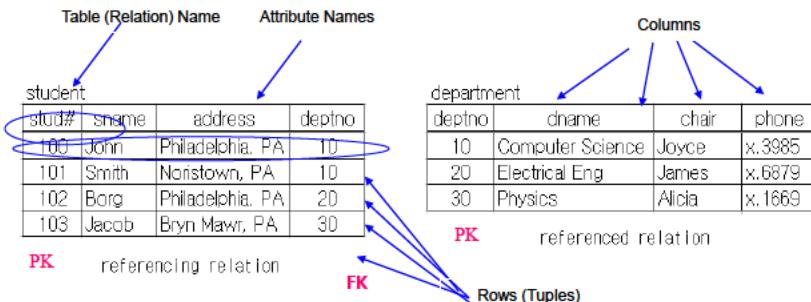
Basics in the Relational Model

- Data representation: a set of tables with structured data
- Primary keys (Candidate key, Alternate keys)
- Foreign Key
- Referential Integrity constraints
- Integrity constraints in RDBMSs
- Null values
- Metadata and Data Dictionary
- Relational schema design using ER diagrams
- Normalization (FD, MVD, JD) and Normal Forms
- SQL (*Aggregation, Subquery, JOIN, PL/SQL functions/procedures/triggers*)
- ACID property in multi-user transactions
- Why relational databases have been popular in industry?

Terminology in RDB

Relational Schema

STUDENT (stud#, sname, address, deptno)
DEPARTMENT (deptno, dname, chair, phone)



A Primary Key (PK) uniquely identifies each row.

A Foreign Key (FK) is a set of attributes
that are primary key values from another table.

7

Characteristics of Relational Model

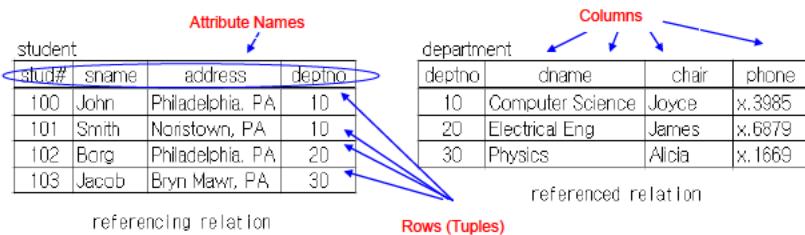


TABLE
3.1 Characteristics of a Relational Table

1	A table is perceived as a two-dimensional structure composed of rows and columns.
2	Each table row (tuple) represents a single entity occurrence within the entity set.
3	Each table column represents an attribute, and each column has a distinct name.
4	Each row/column intersection represents a single data value.
5	All values in a column must conform to the same data format.
6	Each column has a specific range of values known as the attribute domain .
7	The order of the rows and columns is immaterial to the DBMS.
8	Each table must have an attribute or a combination of attributes that uniquely identifies each row.

Example of a DDL in SQL Defining a Relational Schema

student			
stud#	sname	address	deptno
100	John	Philadelphia, PA	10
101	Smith	Norristown, PA	10
102	Borg	Philadelphia, PA	20
103	Jacob	Bryn Mawr, PA	30

department			
deptno	dname	chair	phone
10	Computer Science	Joyce	x3885
20	Electrical Eng	James	x6795
30	Physics	Alice	x1669

referencing relation

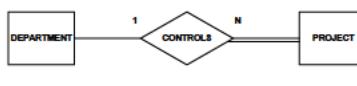
```
CREATE TABLE Department (
    DeptNo      NUMBER(5)  PRIMARY KEY,
    DName       VARCHAR2(15) NOT NULL ,
    Chair        VARCHAR2(20) ,
    Phone        CHAR (10) ,
);
```

```
CREATE TABLE Student (
    Stud#        NUMBER(8) ,
    SName        VARCHAR2(20) NOT NULL ,
    Address      VARCHAR2(40) ,
    DeptNo      NUMBER (5) ,
    CONSTRAINT Stud_PK PRIMARY KEY (Stud#),
    CONSTRAINT Stud_FK FOREIGN KEY (DeptNo)
        REFERENCES Department (DeptNo)
);
```



ER Model: Popular Notations

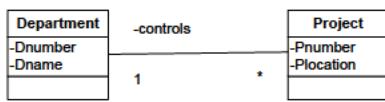
Chen



Visio's Relational notation



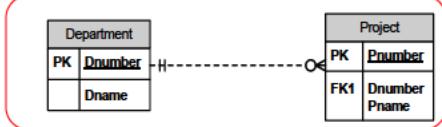
UML



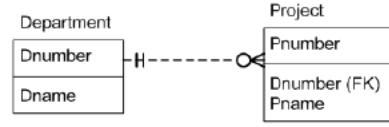
IDEF1X



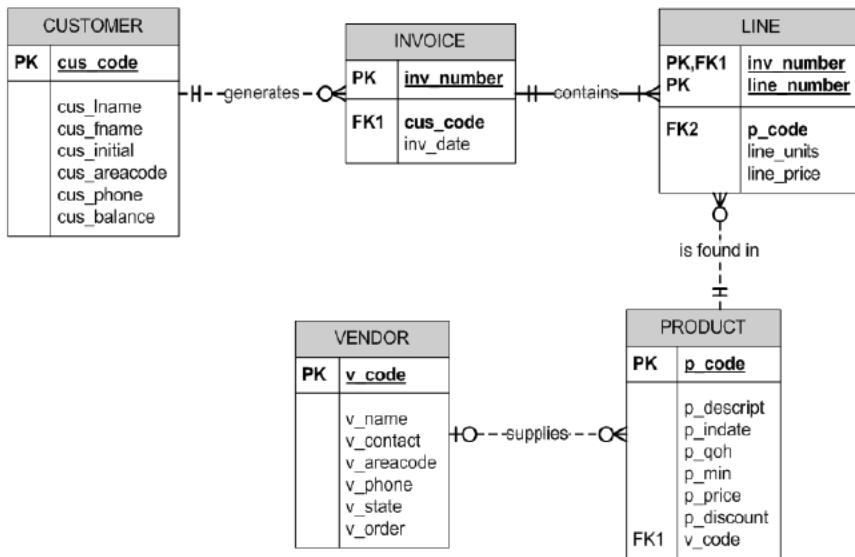
Crow's foot with Relational Notation in Visio



Visio IDEF1X with Crow's foot



ERD for Relational DBs



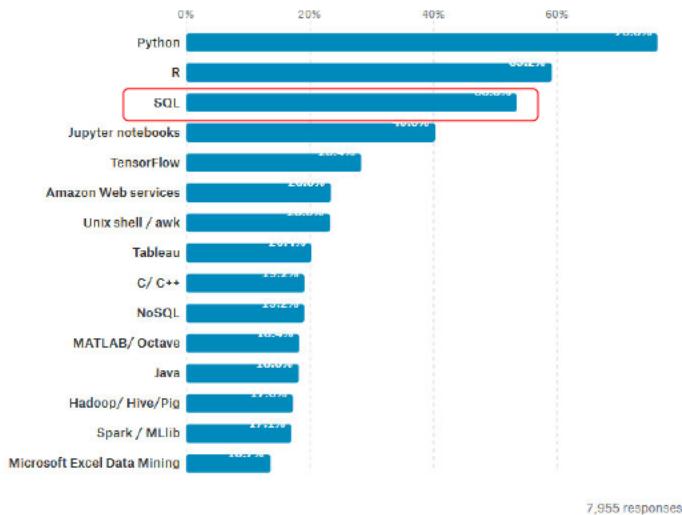
11

11

Why Relational Databases?

- Logical simplicity of the schema: Tables
- Easy, powerful, standard database language: **SQL**
- Transaction reliability: ACID property (Atomicity, Consistency, Isolation, Durability)
- Ad-hoc query processing
- Mature and reliable technologies
- Commercial investment for the last 40 years
- A large group of man-power and user bases

Top Tools used in Data Science (2017)



Source: <https://www.kaggle.com/surveys/2017>



Il-Yeol Song, Ph.D.

13

Top Desired Skills in Data Science (2021)

11 Skills Data Scientists Need

1. Data Visualization
2. Python
3. SQL/NoSQL
4. Social Media Mining
5. Fundamental Statistics
6. Natural Language Processing/Machine Learning
7. Microsoft Excel
8. High-Level Math
9. Teamwork
10. Communication
11. Business Savvy

Source: <https://bootcamp.berkeley.edu/blog/data-scientist-skills/>



Il-Yeol Song, Ph.D.

14

14

Two major Components of SQL

Data Definition Language (DDL)

- CREATE TABLE
- ALTER TABLE
- DROP TABLE
- CREATE/DROP VIEW
- CREATE/DROP INDEX

Data Manipulation Language(DML)

- UPDATE TABLE
- INSERT INTO
- DELETE FROM
- **SELECT** FROM

DDL manages metadata
DML manages data



Il-Yeol Song, Ph.D.

15

A Typical SQL Statement and Rules

STUDENT

ID	Name	EnrollYear

- **SELECT** enrollyear **FROM** student **WHERE** name =‘John’;
- SQL statement comprises
 - **Keywords**
 - Case-insensitive
 - Basic command set: fewer than 100 English words
 - Space-independent
 - **User-defined word (table names, attribute names)**
 - Case-insensitive
 - **Data**
 - **Case-sensitive** (In MS Access, case-insensitive)
 - Ends with a semi-colon



Il-Yeol Song, Ph.D.

16

16

An Overview of SQL Commands

DDL

--Changing attribute type

ALTER TABLE project

```
MODIFY (budget  
NUMBER(9,2));
```

--Adding an attribute

ALTER TABLE project **ADD**
(manager CHAR(10));

-- Delete a column from a table

ALTER TABLE project

```
DROP COLUMN  
manager;
```

--Removing a table from database

```
DROP TABLE project;
```

DML

--Inserting a row to a table

```
INSERT INTO project VALUES (1234,  
'Perfect Project', NULL, 'John');
```

--Changing a value of attribute

```
UPDATE project SET budget =  
1.1*budget WHERE projno > 1000;
```

--Deleting a row from a table

```
DELETE FROM project WHERE  
manager = 'John';
```

Queries using SELECT Statement

SELECT

Specifies which columns are to appear in output

FROM

Specifies table(s) to be used

WHERE

Filters rows with conditions

GROUP BY

Forms groups of rows with same column value.

HAVING

Filters groups subject to some condition.

ORDER BY

Specifies the order of the output.

- *Only SELECT and FROM are mandatory.*
- *Order of the clauses cannot be changed.*

An Overview of Simple Commands

- (1) **SELECT * FROM Instructor;**
- (6) **SELECT * FROM Instructor WHERE deptCode **IS NULL**;**
- (2) **SELECT DISTINCT fName FROM Instructor;**
- (7) **SELECT * FROM Instructor WHERE deptCode **IS NOT NULL**;**
- (3) **SELECT fName, lName, ssn FROM Instructor WHERE deptCode = 'math';**
- (8) **SELECT * FROM Instructor WHERE bonus **BETWEEN** 500 **AND** 1000;**
- (4) **SELECT fName, lName FROM Instructor WHERE bonus > 1000;**
- (9) **SELECT * FROM HR WHERE hireDate **BETWEEN** '01-MAY-2012' **AND** '31-MAY-2012';**
- (5) **SELECT * FROM Instructor WHERE (bonus **>=** 500 **AND** bonus **<=** 1000);**
- (10) **SELECT * FROM Instructor WHERE bonus **IN** (100, 200, 300);**

An Overview of Simple Commands

- (11) **SELECT * FROM Instructor WHERE position = 'assistant';**
- (15) **SELECT fName, lName FROM Instructor WHERE address **LIKE** '%Houston,TX%';**
- (12) **SELECT * FROM R WHERE lName **>** 'S';**
- (16) **SELECT (salary+bonus) FROM Instructor;**
- (13) **SELECT * FROM Instructor WHERE fName **LIKE** 'J%';**
- (17) **SELECT (salary+bonus) **AS** total_income FROM Instructor;**
- (14) **SELECT * FROM Instructor WHERE fName **LIKE** 'J_h_';**
- (18) **SELECT ((A1/3.14)*2.54) **AS** A1_IN_CM FROM R;**
- (19) **SELECT **SYSDATE** AS TODAY FROM DUAL;**
- (20) **SELECT Fname, Lname, Bday, TRUNC(MONTHS_BETWEEN(**SYSDATE**, Bday)/12) **AS** "Actual Age" FROM Person;**

DATE Manipulation

- Show today's date in the form of MM/DD/YYYY format

```
SELECT SYSDATE, TO_CHAR(SYSDATE, 'MM/DD/YYYY') as
      CurrDate FROM DUAL;
```

- Show today's date including time

```
SELECT SYSDATE, TO_CHAR(SYSDATE, 'yyyy-mm-dd hh24:mi:ss')
      as CurrDateTime FROM DUAL;
```

- Find products whose INDATE is more than 90 days old

```
SELECT      P_Code, P_Desc, P_Indate
  FROM      Product
 WHERE      P_Indate <= SYSDATE - 90;
```

- Find all the orders received in Feb 2021.

```
SELECT Order#, Odate
  FROM ORDER
 WHERE Odate BETWEEN '01-Feb-21' AND '28-Feb-21';
```

STRING Manipulation

- **Concatenating strings**

```
SQL> SELECT Lname||', '||Fname "FULL NAME"  FROM Emp;
will display
FULL NAME
Bond, James
```

- **Handling Capitalization**

```
SQL> SELECT Lname, INITCAP(Lname), UPPER(Lname), LOWER(Lname)
  FROM Emp;
```

- **Capitalization:** Show only the first letter in capital and the rest in lower

SQL> SELECT INITCAP(LOWER(Lname)) FROM Emp;

- **Padding blanks to strings**

RPAD: Pad to the right side of the column with left justification
 LPAD: Pad to the left side of the column with right justification
 SQL> SELECT RPAD(Lname, 15, '!'), Age, LPAD(Lname, 15), FROM EMP:

Clinton.....	53Clinton
Gore.....	52Gore

STRING Manipulation

```

/* TRIM function is used to remove all leading or trailing characters (or both)
   from a character string.*/
SQL>SELECT TRIM(" LEADING Fellows ") FROM DUAL;

/* Use of LTRIM(left trim) and RTRIM(right trim) functions*/
SQL> SELECT Lname, LTRIM(Lname, 'SA'), RTRIM(Lname, 'S')
   FROM Emp;
(LTRIM prints M from SAM, RTRIM prints WALE from WALES)

/* Remove . at the end and "THE from the front */
SQL> SELECT LTRIM (RTRIM (Title,'.'), 'THE') FROM Magazine;
Will convert
      MY DARING."
      "THE GOD FATHER."
into
      MY DARING
      GOD FATHER

```



STRING Manipulation

```

/* INSTR(string, set [, start [, occurrence]]) return the position number where in the
   string it found the "set", starting from the start position */

/*Show names and the position number which have 'E' in the 5th or greater position */
SQL> SELECT Lname, INSTR(Lname, 'E', 5) "AFTER FIVE"
   FROM Emp;

/* SUBSTR(String, starting_position, #chars to be returned)
   Extract numeric part, add 1000, and concatenate */
SQL> SELECT Lname, E_ID,
   'S' || TO_CHAR( TO_NUMBER (SUBSTR (E_ID,2,3) ) + 1000) "NEW E_#"
   FROM Emp;

Will display
      LNAME  E_ID          NEW E_#
      Jones   E001          S1001
      Wales   E002          S1002

```



Aggregation

- Example of five aggregation functions
SELECT MAX(bonus), MIN(bonus), AVG(bonus), COUNT(bonus), SUM(bonus)
FROM Instructor
WHERE position = 'assistant';
- Note: What's wrong with the following?
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > AVG(bonus);

We can use aggregate functions only in **SELECT** and **HAVING** clause

- How many different course titles are there?

```
SELECT      COUNT(DISTINCT title) AS count
FROM        Course;
```

Descriptive Statistics in SQL

- **SELECT cus_fname, cus_lname, cus_balance**
AVG(cus_balance) mean,
MEDIAN (cus_balance) median,
STATS_MODE (cus_balance) mode,
STDDEV (cus_balance) "Standard Deviation"
FROM Product;

CUSTOMER	
PK	<u>cus_code</u>
	cus_lname
	cus_fname
	cus_initial
	cus_areacode
	cus_phone
	cus_balance

- Computing median over partition

```
SELECT cus_fname, cus_lname, cus_balance
MEDIAN (cus_balance) OVER (PARTITION BY cus_areacode) AS median
FROM Product;
```

- Computing SD over unique values

```
SELECT cus_fname, cus_lname, cus_balance
STDDEV (DISTINCT cus_balance) FROM Product;
```

Aggregation with GROUP BY

Find the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT deptCode, COUNT(instructorID) AS count, SUM(bonus) AS sum
FROM Instructor
GROUP BY deptCode
ORDER BY deptCode;
```

- Result:

deptCode	count	sum
acct	2	1100
math	2	300

Instructor
-instructorID(PK)
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus

- **GROUP BY must include all non-aggregate function column names in the SELECT list.**

Aggregation with GROUP BY and HAVING

```
+ Oracle SQL*Plus
File Edit Search Options Help
SQL> SELECT v_code, COUNT(DISTINCT (p_code)), AVG(p_price)
  2  FROM product
  3 GROUP BY v_code;
    U_CODE COUNT(DISTINCT(P_CODE)) AVG(P_PRICE)
-----+
    21225                  2      8.47
    21231                  1      8.45
    21344                  3     12.49
    23119                  2     41.97
    24288                  3   155.593333
    25595                  3     89.63
                  2     18.135
7 rows selected.

SQL> SELECT v_code, COUNT(DISTINCT (p_code)), AVG(p_price)
  2  FROM product
  3 GROUP BY v_code
  4 HAVING AVG(p_price) < 10;
    U_CODE COUNT(DISTINCT(P_CODE)) AVG(P_PRICE)
-----+
    21225                  2      8.47
    21231                  1      8.45
```

PRODUCT	
PK	v_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	

Aggregation with GROUP BY and HAVING

- Find the number of invoice per year and per month for the last 3 years

```

SELECT EXTRACT(year FROM inv_date) "Year",
       EXTRACT(month FROM inv_date) "Month",
       COUNT(inv_date) "No. of Invoices"
FROM invoice
GROUP BY EXTRACT(year FROM inv_date),
         EXTRACT(month FROM inv_date)
HAVING EXTRACT(year FROM inv_date) IN (2017, 2016, 2015)
ORDER BY "No. of Invoices" DESC;

```

INVOICE	
PK	inv_number
FK1	cus_code inv_date

CASE Statement

- Simple CASE Statements

```

SELECT      Fname, Lname,
(CASE      DNO
    WHEN 1      THEN 'Headquarters'
    WHEN 4      THEN 'Administration'
    WHEN 5      THEN 'Research'
    ELSE        'No department'
END)      AS Department
FROM Employee;

```

Output:

Fname	Lname	Department
John	Smith	Research
Franklin	Wong	Research
Alica	Zelaya	Administration

CASE Statement

- Searched CASE Statements

```
SELECT Fname, Lname, Salary
  (CASE Salary
    WHEN Salary <= 25000 THEN 1500
    WHEN Salary > 25000 AND Salary < 50000 THEN 1000
    WHEN Salary > 50000 AND Salary < 100000 THEN 500
    ELSE 0
  END) "Bonus"
FROM Employee;
```

Output:

Fname	Lname	Salary	Bonus
John	Smith	30000	1000
Franklin	Wong	40000	1000
Alica	Zelaya	25000	1500

Nested Subquery and Aggregate Functions

- Find the product that has the oldest date

```
SELECT P_CODE, P_DESCRIP
FROM PRODUCT
WHERE P_INDATE = (
  SELECT MIN (P_INDATE)
  FROM PRODUCT);
```

PRODUCT	
PK	p_code
	p_descrip
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

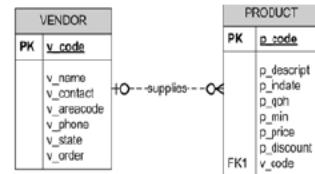
This query is in the form of a nested query.

JOIN: Two Syntax

JOIN is usually done through ***PK-FK chains***

- For each product, find their vendor code and names

```
SELECT      P.P_Code, V.V_Code, V.V_Name
FROM        Product P, Vendor V
WHERE       P.V_Code = V.V_code;
```



```
SELECT      P.P_Code, V.V_Code, V.V_Name
FROM        Product P INNER JOIN Vendor V
ON          P.V_Code = V.V_code;
```

Types of JOINS

TABLE 8.1

SQL JOIN EXPRESSION STYLES

JOIN CLASSIFICATION	JOIN TYPE	SQL SYNTAX EXAMPLE	DESCRIPTION
CROSS	CROSS JOIN	SELECT * FROM T1, T2	Returns the Cartesian product of T1 and T2 (old style)
		SELECT * FROM T1 CROSS JOIN T2	Returns the Cartesian product of T1 and T2
INNER	Old-style JOIN	SELECT * FROM T1, T2 WHERE T1.C1=T2.C1	Returns only the rows that meet the join condition in the WHERE clause (old style); only rows with matching values are selected
	NATURAL JOIN	SELECT * FROM T1 NATURAL JOIN T2	Returns only the rows with matching values in the matching columns; the matching columns must have the same names and similar data types
	JOIN USING	SELECT * FROM T1 JOIN T2 USING (C1)	Returns only the rows with matching values in the columns indicated in the USING clause
	JOIN ON	SELECT * FROM T1 JOIN T2 ON T1.C1=T2.C1	Returns only the rows that meet the join condition indicated in the ON clause
OUTER	LEFT JOIN	SELECT * FROM T1 LEFT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the left table (T1) with unmatched values
	RIGHT JOIN	SELECT * FROM T1 RIGHT OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from the right table (T2) with unmatched values
	FULL JOIN	SELECT * FROM T1 FULL OUTER JOIN T2 ON T1.C1=T2.C1	Returns rows with matching values and includes all rows from both tables (T1 and T2) with unmatched values

Types of JOINS

INNER JOIN



```
SELECT *
FROM A
INNER JOIN B ON A.key = B.key
```

LEFT JOIN



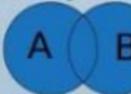
```
SELECT *
FROM A
LEFT JOIN B ON A.key = B.key
```

RIGHT JOIN



```
SELECT *
FROM A
RIGHT JOIN B ON A.key = B.key
```

FULL JOIN



```
SELECT *
FROM A
FULL JOIN B ON A.key = B.key
```

NATURAL JOIN:

```
SELECT *
FROM table_A
NATURAL JOIN table_B;
```

SELF JOIN:

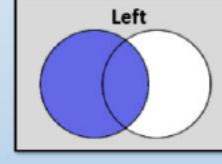
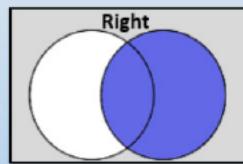
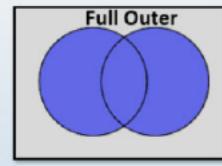
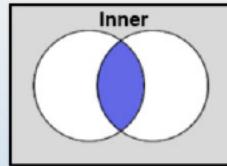
```
SELECT *
FROM table_A X, table_A Y
WHERE X.A = Y.A;
```

CROSS JOIN:

```
SELECT *
FROM table_A
CROSS JOIN table_B;
```

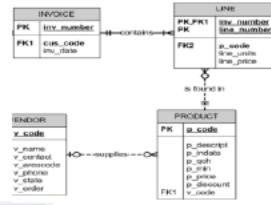
Types of JOINS in Tableau


Tableau
Joins



INNER JOIN

- For each invoice, find the product description, #units, and price of items in the invoice.



SQL> SELECT INVOICE.INV_NUMBER, PRODUCT.P_CODE, P_DESCRPT, LINE_UNITS, LINE_PRICE
2 FROM INVOICE JOIN LINE ON INVOICE.INV_NUMBER = LINE.INV_NUMBER
3 JOIN PRODUCT ON LINE.P_CODE = PRODUCT.P_CODE;

INV_NUMBER	P_CODE	P_DESCRPT	LINE_UNITS	LINE_PRICE
1001	13-Q2/P2	7.25-in. pwr. saw blade	1	14.99
1001	23109-HB	Claw hammer	1	9.95
1002	54778-2T	Rat-tail file, 1/8-in. fine	2	4.99
1003	2238/QPD	B&D cordless drill, 1/2-in.	1	38.95
1003	1546-0Q2	Hrd. cloth, 1/4-in., 2x50	1	39.95
1003	13-Q2/P2	7.25-in. pwr. saw blade	5	14.99
1004	54778-2T	Rat-tail file, 1/8-in. fine	3	4.99
1004	23109-HB	Claw hammer	2	9.95
1005	PVC23DRT	PVC pipe, 3.5-in., 8-ft	12	5.87
1006	SM-18Z77	1.25-in. metal screw, 25	3	6.99
1006	2232/OTY	B&D jigsaw, 12-in. blade	1	109.92
1006	23109-HB	Claw hammer	1	9.95
1006	89-WRE-Q	Hicut chain saw, 16 in.	1	256.99
1007	13-Q2/P2	7.25-in. pwr. saw blade	2	14.99
1007	54778-2T	Rat-tail file, 1/8-in. fine	1	4.99
1008	PVC23DRT	PVC pipe, 3.5-in., 8-ft	5	5.87
1008	WR3/TT3	Steel matting, 4'x8'x1/6", .5" mesh	3	119.95
1008	23109-HB	Claw hammer	1	9.95

18 rows selected.

SQL> -



(Coronel & Morris, 2016)

37

LEFT OUTER JOIN

- For ALL the Vendors, find all the products they provide. Include all the vendors whether or not they provide any product.



SQL> SELECT P_CODE, VENDOR.V_CODE, V_NAME
2 FROM VENDOR LEFT JOIN PRODUCT ON VENDOR.V_CODE = PRODUCT.V_CODE;

P_CODE	V_CODE	V_NAME
11QER/31	25595	Rubicon Systems
13-Q2/P2	21344	Gomez Bros.
14-Q1/L3	21344	Gomez Bros.
1546-QQ2	23119	Randsets Ltd.
1558-QW1	23119	Randsets Ltd.
2232/OTY	24288	ORDVA, Inc.
2232/QWE	24288	ORDVA, Inc.
2238/QPD	25595	Rubicon Systems
23109-HB	21225	Bryson, Inc.
54778-2T	21344	Gomez Bros.
89-WRE-Q	24288	ORDVA, Inc.
SM-18Z77	21225	Bryson, Inc.
SW-23116	21231	D&E Supply
WR3/TT3	25595	Rubicon Systems
	22567	Dome Supply
	21226	SuperLoo, Inc.
	24004	Brackman Bros.
	25501	Damat Supplies
	25443	B&K, Inc.

19 rows selected.

SQL> -

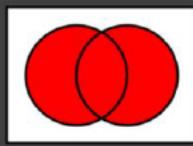


(Coronel & Morris, 2016)

38

SET Operations

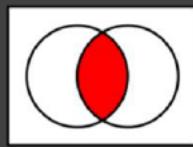
```
-- All customers and staff
SELECT first_name, last_name
FROM customer
UNION
SELECT first_name, last_name
FROM staff
```



UNION

R ∪ S

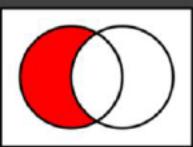
```
-- customers that are actors
SELECT first_name, last_name
FROM customer
INTERSECT
SELECT first_name, last_name
FROM actor
```



INTERSECT

R ∩ S

```
-- customers that are not staff
SELECT first_name, last_name
FROM customer
EXCEPT
SELECT first_name, last_name
FROM staff
```



MINUS

R - S



Il-Yeol Song, Ph.D.

39

39

SQL Analytic Functions: ROLLUP

- Example: **GROUP BY ROLLUP (V_Code, P_Code)**
produces the union of
 - GROUP BY V_Code, P_Code
 - GROUP BY V_Code
 - Grand total

```
SQL> SELECT    V_CODE, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
  2  FROM      DWDSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
  3  GROUP BY ROLLUP (V_CODE, P_CODE)
  4  ORDER BY V_CODE, P_CODE;
```



Il-Yeol Song, Ph.D.

40

40

Oracle SQL*Plus

```
File Edit Search Options Help
SQL> SELECT V_CODE, P_CODE, SUM(TRADE_UNITS*TRADE_PRICE) AS TOTSALES
  2 FROM DWTRASALEFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
  3 GROUP BY ROLLUP (V_CODE, P_CODE)
  4 ORDER BY V_CODE, P_CODE;
```

V_CODE	P_CODE	TOTSALES
21225	23109-HB	99.5
21225	PUC23DRT	199.58
21225	SM-18277	41.94
21225		341.02
21344	13-Q2/P2	239.84
21344	54778-2T	59.88
21344		299.72
23119	1546-002	79.9
23119		79.9
24288	2232/QTY	219.84
24288	89-WRE-Q	513.98
24288		733.82
25595	2238/QPD	77.9
25595	WR3/TT3	719.7
25595		797.6
		2252.06

Subtotals by V_CODE

Grand total for all P_CODE values

16 rows selected.

(Coronel & Morris, 2016)

41

SQL Analytic Functions: Cube

- The CUBE extension
 - Enables you to get a subtotal for each column listed in the expression, in addition to a grand total for the last column listed
 - GROUP BY CUBE (TM_Month, P_CODE) is a union of:
 - GROUP BY CUBE (TM_Month, P_CODE)
 - GROUP BY CUBE (TM_Month)
 - GROUP BY CUBE (P_CODE)
 - Grand total

```
SQL> SELECT TM_MONTH, P_CODE, SUM(TRADE_UNITS*TRADE_PRICE) AS TOTSALES
  2 FROM DWTRASALEFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWTIME
  3 GROUP BY CUBE (TM_MONTH, P_CODE)
  4 ORDER BY TM_MONTH, P_CODE;
```

Oracle SQL*Plus

```
SQL> SELECT TH_MONTH, P_CODE, SUM(SALE_UNITS*SALE_PRICE) AS TOTSALES
  2  FROM DW_DAYSALESFACT NATURAL JOIN DW_PRODUCT NATURAL JOIN DW_TIME
  3 GROUP BY CUBE (TH_MONTH, P_CODE)
  4 ORDER BY TH_MONTH, P_CODE;
```

TH_MONTH	P_CODE	TOTSALES
9	13-Q2/P2	134.91
9	1546-QQ2	79.9
9	2232/QTY	109.92
9	2238/QPD	77.9
9	23109-HB	59.7
9	54778-2T	39.92
9	89-WRE-Q	256.99
9	PUC23DRT	99.79
9	SH-18277	28.97
9	WR3/TT3	359.85
9		1239.85
10	13-Q2/P2	104.93
10	2232/QTY	109.92
10	23109-HB	39.8
10	54778-2T	19.96
10	89-WRE-Q	256.99
10	PUC23DRT	99.79
10	SH-18277	28.97
10	WR3/TT3	359.85
10		1012.21
11	13-Q2/P2	239.84
11	1546-QQ2	79.9
11	2232/QTY	219.84
11	2238/QPD	77.9
11	23109-HB	99.5
11	54778-2T	59.88
11	89-WRE-Q	513.98
11	PUC23DRT	199.58
11	SH-18277	41.94
11	WR3/TT3	719.7
11		2252.06

(Coronel & Morris, 2016)

31 rows selected.

Subtotals by month

Subtotals by product

Grand total for all products and months

43

RANK and DENSE_RANK Functions

- RANK() leave ranking gaps when there multiple rows in the same rank (1, 2, 2, 4, 5...)
- DENSE_RANK() does **not** leave ranking gaps (1, 2, 2, 3, 4, 5)

```
SELECT p_code, p_descript, p_price,
       RANK() OVER
          (ORDER BY p_price NULLS LAST) AS Rank,
       DENSE_RANK()
          (ORDER BY p_price NULLS LAST) AS Dense_rank
FROM   Product
ORDER BY p_price;
```

PRODUCT	
PK	p_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

- DESC means descending order
- NULLS LAST means null values are smaller than non-null values
- You may use NULLS FIRST

Il-Yeol Song, Ph.D.

44

44

RANK and DENSE RANK Functions

- Ranking applied to dates

```
SELECT p_code, p_descript, p_indate,
       RANK() OVER
           (ORDER BY p_indate NULLS LAST) AS Rank,
       DENSE_RANK()
           (ORDER BY p_indate NULLS LAST) AS Dense_rank
  FROM Product
 ORDER BY p_indate;
```

PRODUCT	
PK	p_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Top-N Queries

- Before Oracle 12c:

```
SELECT p_code, p_descript,
       (SELECT p_price,
               RANK() OVER
                   (ORDER BY p_price NULLS LAST) AS Rank
      FROM product)
  FROM Product
 WHERE Rank <= 10
 ORDER BY p_price;
```

- In Oracle 12c:

```
SELECT p_code, p_descript, p_price,
       RANK() OVER
           (ORDER BY p_price NULLS LAST) AS Rank,
  FROM Product
 ORDER BY p_price;
 FETCH FIRST 10 ROWS ONLY;
```

- Another option by percentage
FETCH FIRST 20 PERCENT ROWS ONLY;
- MySQL uses LIMIT clause
LIMIT 10

RANK with PARTITION BY

```
SELECT employee_id, first_name, department_id, salary,  
RANK() OVER (PARTITION BY department_id  
    ORDER BY salary) rank
```

```
FROM employee;
```

	EMPLOYEE_ID	FIRST_NAME	DEPARTMENT_ID	SALARY	RANK
1	119	Karen	30	2500	1
2	118	Guy	30	2600	2
3	117	Sigal	30	2800	3
4	116	Shelli	30	2900	4
5	115	Alexander	30	3100	5
6	114	Den	30	11000	6
7	144	Peter	50	2500	1
8	143	Randall	50	2600	2
9	142	Curtis	50	3100	3
10	141	Trenna	50	3500	4
11	107	Diana	60	4200	1
12	105	David	60	4800	2



Il-Yeol Song, PhD.

47

47

Windows Function: WIDTH_BUCKET

- This operation assigns a bucket number to the expression that it evaluates. It helps in generating equi-width histograms.
 - Example: Show the last names, salaries and bucket numbers of all employees according to the following rule:
 - Salary <30,000 -> 0 30,000-40,000 -> 1
 - 40,000-50,000 -> 2 Salary >50,000 -> 3

SQL> SELECT lname, salary,

WIDTH_BUCKET(salary, 30000, 50000, 2) AS Bucket

FROM employee;

LNAME	SALARY	BUCKET
Smith	30000	1
Wong	40000	2
Zelaya	25000	0
Wallace	43000	2
Narayan	38000	1
English	25000	0
Jabbar	25000	0
Borg	55000	3

8 rows selected.



Il-Yeol Song, Ph.D.

48

48

Other Advanced SQL

- SQL with JSON
- SQL with Python
- Views and Materialized views
- PL/SQL
 - Stored functions
 - Stored Procedures
 - Triggers
 - Packages

Question?

