

Text Processing / Mining / Search

Mat Kelly, PhD

Assistant Professor, Information Science
Drexel University, College of Computing & Informatics (CCI)
Philadelphia, PA

mkelly@drexel.edu
<https://matkelly.com>
[@machawk1](https://twitter.com/machawk1)

LIS Education And Data Science Integrated Network Group (LEADING) Bootcamp
Week 1: June 7, 2023

What We Will Cover

- Text Processing
 - Normalization
 - Regular Expressions (RegEx)
 - Stop words
 - Stemming
 - n-grams
- Text Mining ↑ incorporated into above
- Text Search
 - IR methods
 - Evaluating similarity

Goals of this Presentation

- Introduce aforementioned topics at a high level
- Practical:
 - Process data (text)
 - Be able to identify and extract relevant information
 - Small corpus, intent is to extrapolate
- Basic Information Retrieval concepts
 - Applicable both to text corpora and other kinds.

My Background/Specialization

- Computer Science PhD, research in web archiving
- In Information Science Dept. @ Drexel CCI
- Teaching:
 - INFO624: Information Retrieval
 - INFO600: Web Systems & Architecture
 - INFO202: Data Curation
- Authored multiple Python-based tools while a grad student
 - <https://github.com/machawk1>
- Coming from a Systems Development Background (cf. data science)
 - The tools to generate, consume the data



Sample Corpus (small collection of “Documents”)

- D_1 : The quick brown fox jumped over the lazy dog.
- D_2 : The sum of the square of the legs of a right triangle are equal to the square of the hypotenuse.
- D_3 : We are what we pretend to be, so we must be careful about what we pretend to be.



Text Processing

Preparation

- Tokenize
- Remove Stopwords
- Stem Words
- Construct N-grams

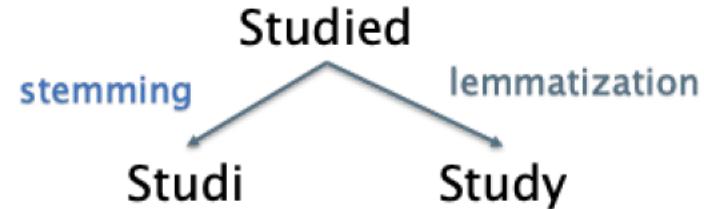
Tokenization



- Counting instances
 - D₃ : We are what we pretend to be, so we must be careful about what we pretend to be.
 - (1) We (2) are (3) what (4) we (5) pretend (6) to (7) be (8) so (9) must (10) careful (11) about
- Normalization
 - dog, dogs → dog
 - leading, leads, leader, leaders → lead
 - Simpler examples, e.g., Leading, leading, LEADING → leading
 - co-education? San Francisco? Hewlett-Packard

Tokenization

- Tokenization, normalization, and inverted index
- Normalization: casefolding, stemming, lemmatization
 - Leading, LEADING, leading → leading
 - C.A.T., cat, Cat → cat?
 - General Motors / general motors?
- Lemmatization: reducing to canonical/dictionary form
 - am, are, is → be
 - car , cars, car's, cars' → car
 - For example:
 - the boy's cars are different colors
 - The boy car be different color



Remove Stopwords

- The quick brown fox jumped over the lazy dog.
- The sum of the square of the legs of a right triangle are equal to the square of the hypotenuse.
- We are what we pretend to be, so we must be careful about what we pretend to be.
- Some bases/rationale for stopwords
 - Common / noisy / not semantically relevant
- To Be Or Not To Be (Hamlet)

Remove Stopwords

- The quick brown fox jumped over the lazy dog.
- The sum of the square of the legs of a right triangle are equal to the square of the hypotenuse.
- We are what we pretend to be, so we must be careful about what we pretend to be.
- Some bases/rationale for stopwords
 - Common / noisy / not semantically relevant
- To Be Or Not To Be (Hamlet)

Remove Stopwords

- quick brown fox jumped over lazy dog.
- sum square the legs a right triangle are equal square hypotenuse.
- We are what we pretend, so we must be careful about what we pretend.

- Some bases/rationale for stopwords
 - Common / noisy / not semantically relevant
- **To Be Or Not To Be** (Hamlet)



Stem Words

- Words “leading”, “leads”, “leaders”, “leader” all have same “root”
 - We want to coalesce for semantic/syntactic analysis
 - **Porter stemmer**: fast but not very precise
 - **Snowball (Porter2)**: more aggressive but generally better than Porter
 - **Lancaster**: most aggressive, least readable result



Typical **Porter** rules

- sses → ss
- ies → i
- ational → ate
- tional → tion

Weight of word sensitive to rules
(m>1) EMENT →

- replacement → replac
- cement → cement

Stem Words

% python3 (code to the right)

D0 stemmed terms: the quick brown fox jump over the lazi dog

.

D1 stemmed terms: the sum of the squar of the leg of a right triangl are equal to the squar of the hypotenuse .

D2 stemmed terms: we are what we pretend to be , so we must be care about what we pretend to be .

```
import nltk
from nltk.stem.porter import *

d1 = "The quick brown fox jumped over the lazy dog."
d2 = "The sum of the square of the legs of a right triangle are equal to the square
      of the hypotenuse."
d3 = "We are what we pretend to be, so we must be careful about what we pretend to
      be."
docs = [d1, d2, d3]

porter_stemmer = PorterStemmer()

term_list = []
# Tokenize each document
for document in docs:
    term_list.append(nltk.word_tokenize(document))

stemmed_term_list = []
# Stem all tokens in each document
for i, term_list_for_doc in enumerate(term_list):
    stemmed_term_list.append([porter_stemmer.stem(word) for word in
                           term_list_for_doc])

# Print out stemmed term list per document
for i, doc in enumerate(docs):
    print(f'D{i} stemmed terms: {" ".join(stemmed_term_list[i])}'")
```

colab



Construct N-grams



- Helps to identify multi-word semantics
- Bi-grams examples:
 - the quick; quick brown; brown fox; fox jumped; jumped over; over the; the lazy; lazy dog
- Tri-gram examples:
 - the quick brown; quick brown fox; brown fox jumped; fox jumped over; etc.
- Useful for identifying compound tokens
 - Drexel University
 - data science
 - The Queen of England

```
import nltk
from nltk.tokenize import word_tokenize

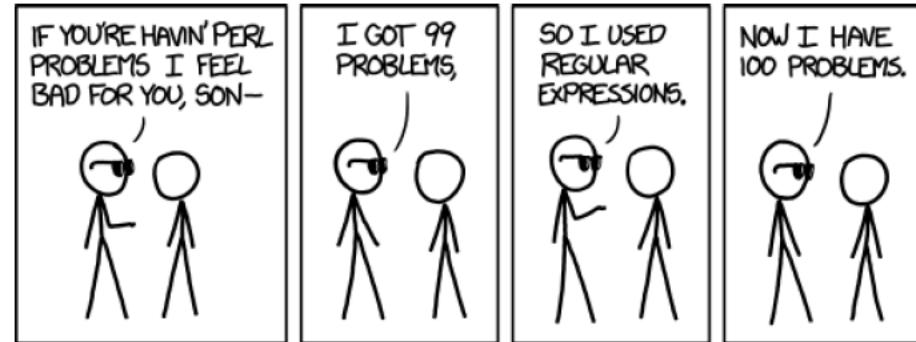
doc = "The quick brown fox jumped over the lazy
dog."
tokens = nltk.word_tokenize(doc)
bigrm = nltk.bigrams(tokens)
print(*map(' '.join, bigrm), sep=', ')
```

> The quick, quick brown, brown fox, fox jumped,
jumped over, over the, the lazy, lazy dog, dog .

colab

Regular Expressions

- Test whether a string matches a pattern
- Useful for search/replace character string
- Commonly implemented: Perl, Python, text editors, UNIX command-line
- Powerful but confusing syntax

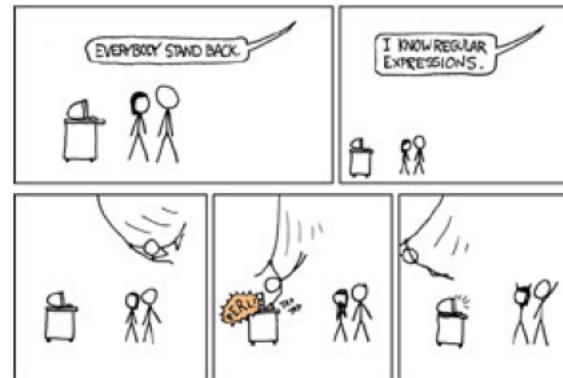


[CC-BY-NC](#) by [XKCD](#)

Regular Expressions

- Basic usage

- Boolean: `analy(s|z)e` → matches analyze, analyse
- Zero-or-one: `colou?r` → matches color, colour
- Zero-or-more: `[a-z]*` → matches a, leading, dog, (empty string), etc. (but not LEADING)
- One-or-more: `b[oO]+` → matches bo, bOooOOOoo, bOO
- Matches n times: `b[o]{2}` → matches boo
- Matches minimally i times: `a[r]{2,}gh` → matches arrgh, arrrrrrrrrrgh
- Matches maximally j times: `b[o]{,4}` → matches boooo, boo, b
- Matches between between i and j times (inclusive): `[bo]{1,3}` → bobo, bo



[CC-BY-NC](#) by [XKCD](#)

How are RegEx useful?

- `re.findall('[a-z]+ of the [a-z]+', d2)`
 - Result: ['sum of the square', 'square of the hypotenuse']
- `re.findall('[a-z]{6,}', d2)`
 - Result: ['square', 'triangle', 'square', 'hypotenuse']
- Can also be compiled for reuse
 - `r = re.compile('\sthe\s[a-z]+')` # find all instance of ' the XXXX'
 - `r.findall(d2)`
 - Result: [' the square', ' the legs', ' the square', ' the hypotenuse']

Shorthands
\s → whitespace
\d → digits
\w → word
...many other

D₂ : The sum of the square of the legs of a right triangle are equal to the square of the hypotenuse.

colab

Search: Evaluation

- Supposing we are searching for text in a collection, how do we measure what is a “good” result?
- An information retrieval standpoint
 - TF-IDF (term frequency - inverse document frequency)
 - Cosine Similarity
 - Precision & Recall

TF-IDF

- Term Frequency ($tf_{t,d}$): The number of times term t appears in document d
- A document with **10** occurrences of the term is *more relevant* than a document with 1 occurrence of the term.
 - But not **10 times** more relevant.
 - We use logarithms to mitigate this effect.
- $tf = \frac{\text{term count}}{\# \text{ of terms in document}}$



Karen Spärck
Jones

By University of
Cambridge
[CC-BY-SA](#)

TF-IDF

- Document Frequency (df)
 - Rare terms more informative than common terms (e.g., stopwords)
 - Give rare terms high weights
- $df = \frac{\text{\# documents with term in it}}{\text{\# of documents}}$
- Inverse df = $\frac{1}{df} = \frac{\text{\# documents}}{\text{\# documents w/ term}}$

higher weights

lower weights

rare terms
(small DF)

frequent terms
(small DF)



Karen Spärck
Jones

By University of
Cambridge
[CC-BY-SA](#)

TF-IDF

- TF-IDF: weight of each term in a document to signify importance
- $\text{tf.idf} = \log(\text{tf} \times \text{idf})$
- Ranking of a document for a query (which is most relevant):

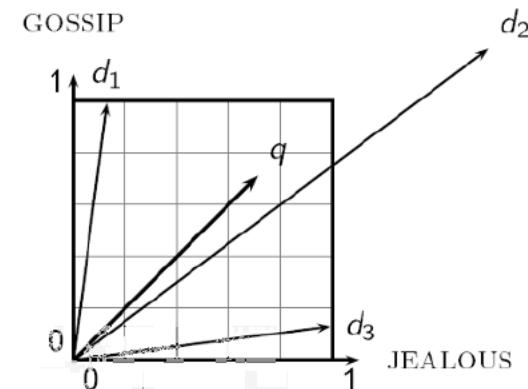
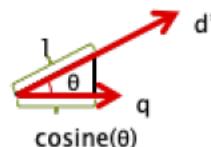
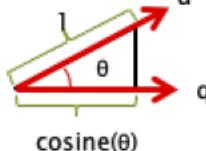
$$\text{Score}(q,d) = \sum_{t \in q \cap d} \text{tf.idf}_{t,d}$$

$$\text{tf.idf} = \frac{\text{term count}}{\# \text{ of terms in document}} \times \frac{\# \text{ documents}}{\# \text{ documents w/ term}}$$

$$\text{tf.idf} = \text{term frequency (tf)} \times \text{inverse document frequency (idf)}$$

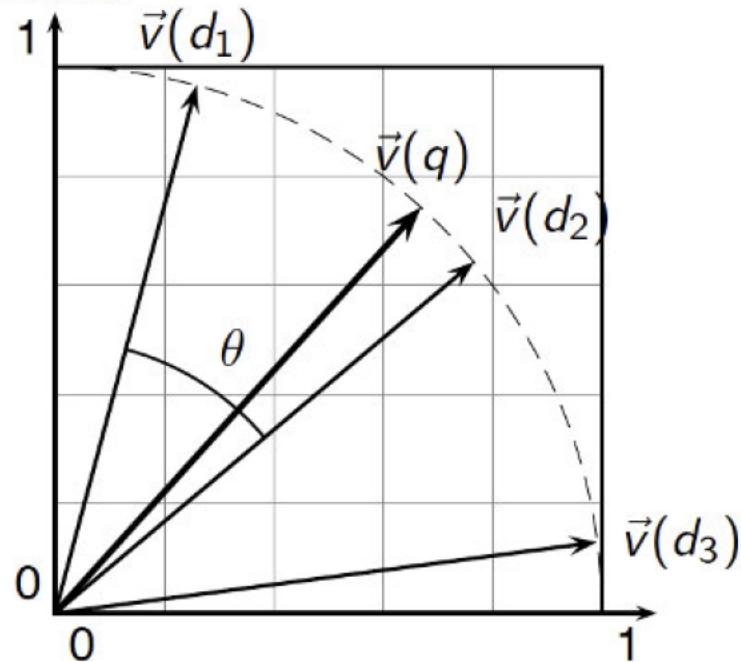
Cosine Similarity

- Plot term vectors, vector length is affected by TF
 - e.g., document with just “be be be be be be”
- Distance between query and document is problematic
- Rather than vector magnitude, use angle between
 - Smaller angle → more similar documents
- To measure how close/similar they are:
 - The Euclidean distance between the two documents can be quite large
 - The angle between the two documents is 0, corresponding to maximal similarity. (Cosine = 1)
- Key idea: Rank documents according to angle with query.

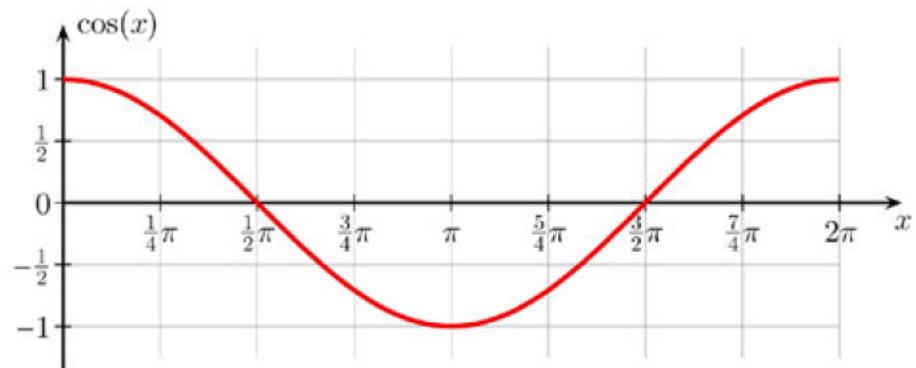


Cosine Similarity Illustrated

POOR



RICH



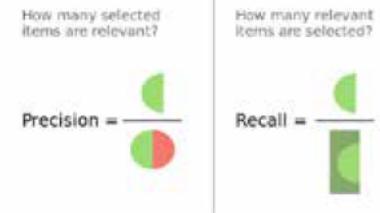
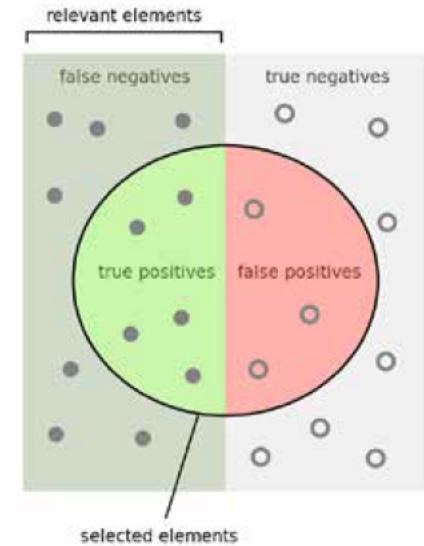
Vector Space Ranking

Using TF-IDF and Cosine Similarity together

- Represent the query as weighted tf-idf vector
- Represent each document as a weighted tf-idf vector
- Compute the cosine similarity score for the query vector and each document vector
- Rank documents with respect to the query by score
- Return top K (e.g., K=10) to the user

Precision & Recall

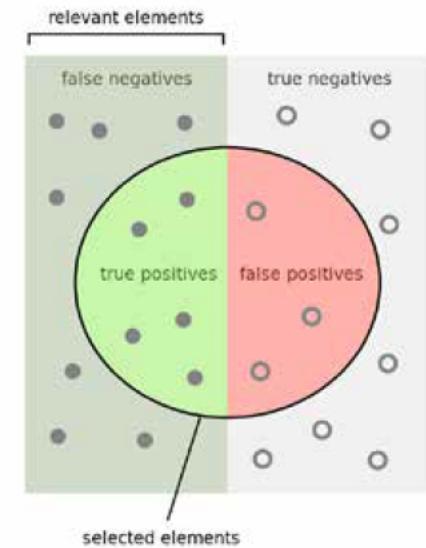
- **True Positive**: Things we got and should have gotten
- **False Positive**: Things we got and should not have gotten
- **False Negative**: Things we did not get and should have got
- **True Negative**: Things we did not get and should not have negative)



[CC-BY-SA](#) by Walbur

Precision & Recall

- Precision & Recall are a means of evaluating retrieval quality (each on a scale of 0.0 to 1.0)
 - What this entails may be domain-specific
 - Generally are inversely proportional
- Precision:** # True positives over # results returned:
 - $TP/(TP+FP)$
- Recall:** True positives over all positives
 - $TP/(TP+FN)$
- Neither metric is sufficient alone, can be gamed
 - High precision/ low recall: very few, but correct result (**P**: 1.0)
 - Low precision / high recall: return everything (**R**: 1.0)



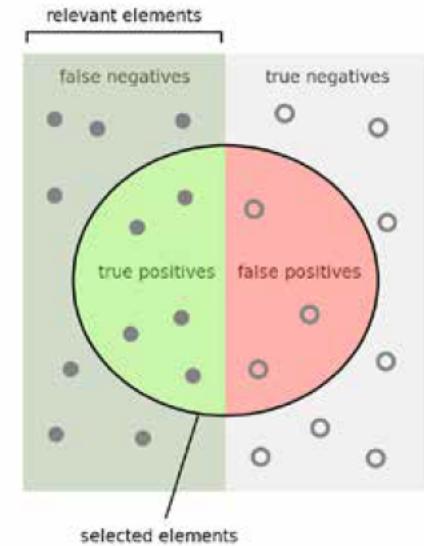
$$\text{Precision} = \frac{\text{How many relevant items are selected?}}{\text{How many selected items are relevant?}}$$

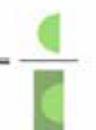
$$\text{Recall} = \frac{\text{How many relevant items are selected?}}{\text{How many relevant items are there?}}$$

[CC-BY-SA](#) by Walbur

Precision & Recall

- Key is to find best balance between **Precision/Recall**
- F-Measure:
 - $$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$
- Measures *accuracy* of the retrieval/classification method
- Range from 0.0 (e.g., recall=1.0) to 1.0 (balance **P/R**)

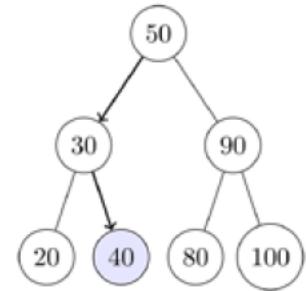


$\text{Precision} = \frac{\text{How many selected items are relevant?}}{\text{How many selected items?}}$ 	$\text{Recall} = \frac{\text{How many relevant items are selected?}}{\text{How many relevant items?}}$ 
---	--

[CC-BY-SA](#) by Walbur

Time/Space Complexity

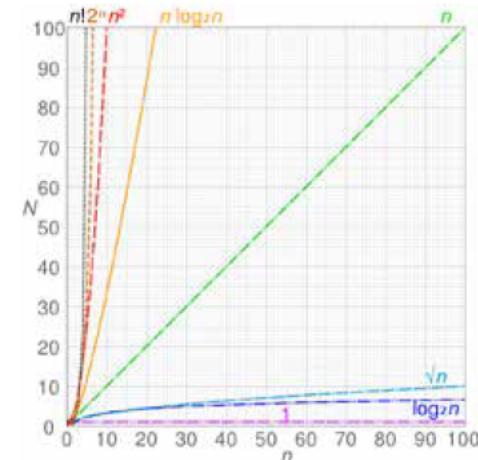
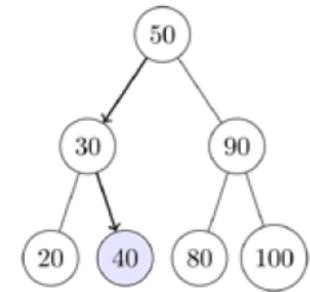
- Big-O notation
 - Search algorithms
 - Sorting algorithms
- Large data sets require efficient processing
- Sorting preferred, makes searching more efficient



[CC BY-SA](#) Cmglee

Time/Space Complexity

- Big-O notation
 - Search algorithms
 - Sorting algorithms
- Large data sets require efficient processing
- Sorting preferred, makes searching more efficient
- **O(n)** → linear → One pass, max look through all

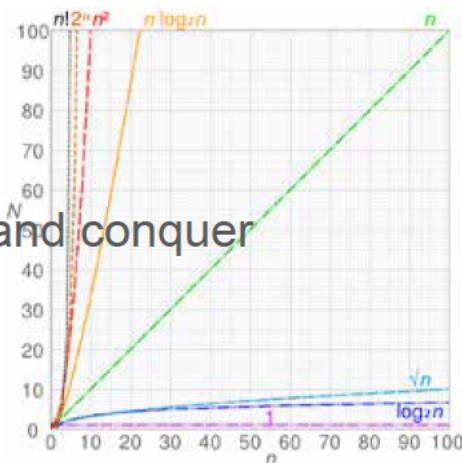
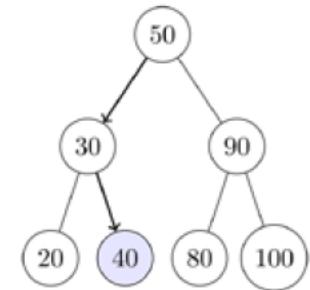


Finding an item in an ordered list

CC BY-SA Cmglee

Time/Space Complexity

- Big-O notation
 - Search algorithms
 - Sorting algorithms
- Large data sets require efficient processing
- Sorting preferred, makes searching more efficient
- $O(n)$ → linear → One pass, max look through all
- **$O(\log n)$** → log → More efficient than linear, requires divide and conquer

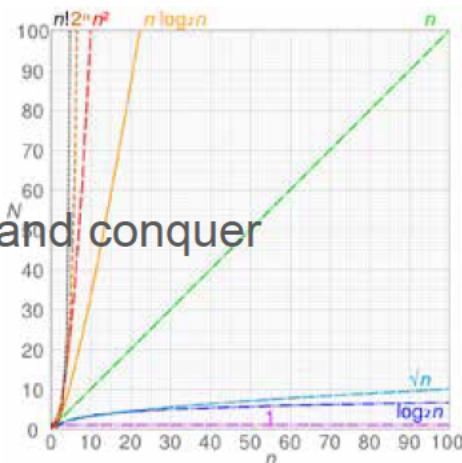
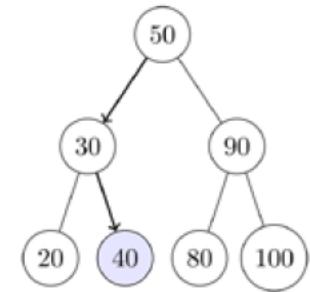


Finding an item a balanced tree structure

CC BY-SA Cmglee

Time/Space Complexity

- Big-O notation
 - Search algorithms
 - Sorting algorithms
- Large data sets require efficient processing
- Sorting preferred, makes searching more efficient
- $O(n)$ → linear → One pass, max look through all
- $O(\log n)$ → log → More efficient than linear, requires divide and conquer
- **$O(n^k)$** → polynomial → Very inefficient, e.g., nested loops



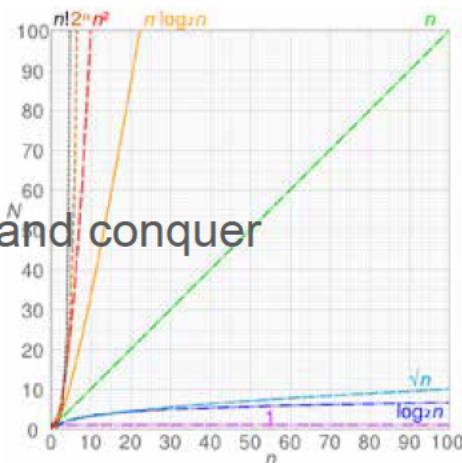
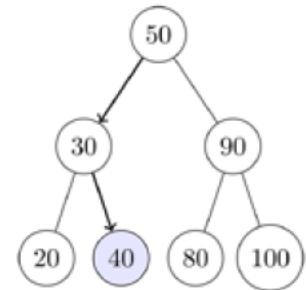
Nested loops

CC BY-SA Cmglee

Time/Space Complexity

- Big-O notation
 - Search algorithms
 - Sorting algorithms
- Large data sets require efficient processing
- Sorting preferred, makes searching more efficient
- $O(n)$ → linear → One pass, max look through all
- $O(\log n)$ → log → More efficient than linear, requires divide and conquer
- $O(n^k)$ → polynomial → Very inefficient, e.g., nested loops
- **$O(n \log n)$** → $n \cdot \log n$ → More efficient than polynomial, less than linear

Various sorting algorithms: Heap, Merge, Quick



CC BY-SA Cmglee

Parallelization

- Means of optimizing execution
- Multiprocessing
 - Execute tasks simultaneously using multiple processors.
 - `multiprocessing.Pool` and `multiprocessing.Process`
- Asynchronous vs. Synchronous execution
 - Sync: retains order, uses locking
 - Async: mixed order, but may finish quicker



Parallelization

- Identifying bottlenecks, programmatic dependencies
 - e.g., corpus metrics cannot be calculated until processing completed on all documents
- Works well for processes that can be chunked or sub-processes executed independently
 - e.g., classification, numerical calculation, relevant text extraction

In closing, topics covered

- Normalization
- Regular expressions
- Tokenization
- Stopwords
- Stemming
- N-grams
- TF-IDF
- Cosine Similarity
- Precision/Recall
- Complexity
- Parallelization