



LEADS 3-Day Camp

Week 2: Data Management:

SQL and NoSQL

Il-Yeol Song, Ph.D.

Professor

College of Computing & Informatics

Drexel University

Philadelphia, PA 19104

Song@drexel.edu

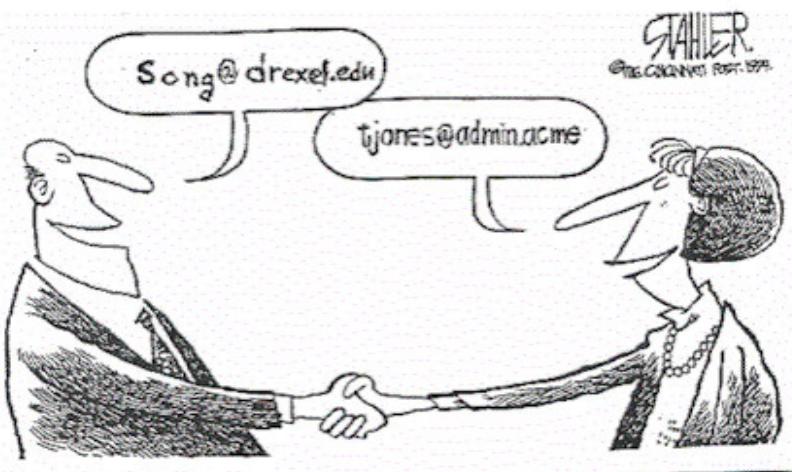
<http://www.cci.drexel.edu/faculty/song/>

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

1

1

Hello!



2

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

2

Instructor

- Professor, College of Computing & Informatics, Drexel University,
- PhD in CS, LSU, Baton Rouge, USA, 1988
- CCI PhD Program Director (2010-2015)
- CVDI Deputy Director (2012-2014)
- Research Topics: *Conceptual Modeling, Data Warehousing, Big Data Management and Data Analytics, Smart Aging*
- Elected as an **ER Fellow**, 2012
- Named an **ACM Distinguished Scientist** in 2013
- Received **Peter Chen Award in Conceptual Modeling** in 2015
- Four teaching awards from Drexel (1991, 2000, 2001, 2011) including **Lindback Distinguished Teaching Award** (2001)
- **Co-Editor-in-Chief**, *Journal of Computing Science & Engineering*
- **Consulting Editor**, *Data & Knowledge Engineering*
- **Co-Chair**, iSchool Model Data Science Curriculum Committee
- **Chair**, IEEE Big Data and Smart Computing (BigComp) Conference
- Published about 200+ papers



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

3

Acknowledgement

- Materials from:
 - Lecture Notes from Drexel CCI INFO 605, 606 and 607 Classes (Il-Yeol Song)
 - Some diagrams and examples from
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management, 13th Edition, 2017. Cengage Learning.



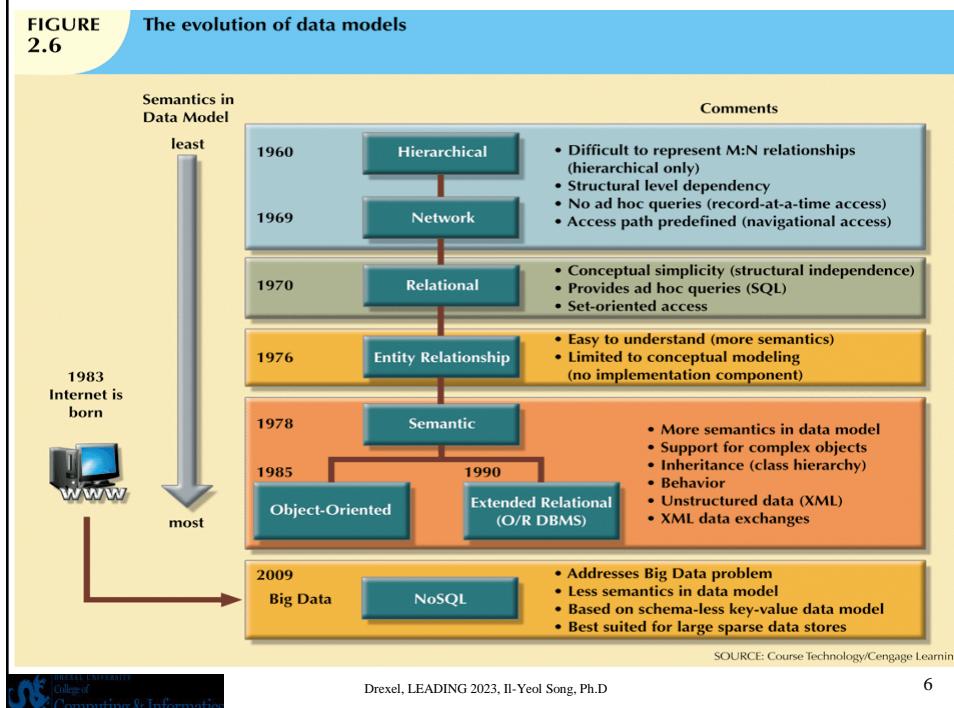
Drexel, LEADING 2023, Il-Yeol Song, Ph.D

4

4

Contents

- Evolution of Data Models
- Relational Databases
- SQL
 - Basic SQL commands
 - SQL for Data Engineering
 - SQL for Analytics
- NoSQL Databases
- MongoDB Basic CRUD Operations



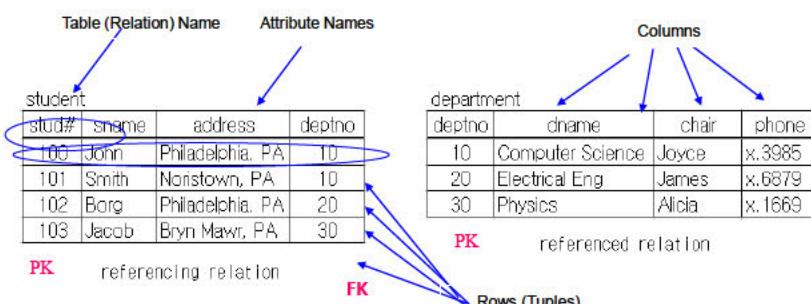
Basics in the Relational Model

- Data representation: a set of tables with structured data
- Primary keys (Candidate key, Alternate keys)
- Foreign Key
- Referential Integrity constraints
- Integrity constraints in RDBMSs (PK, RI, CHECK, Triggers)
- Null values
- Indexes
- Metadata and Data Dictionary
- Relational schema design using ER diagrams
- Normalization (FD, MVD, JD) and Normal Forms
- SQL (DDL and DML: *Aggregation, Subquery, JOIN, PL/SQL functions/procedures/triggers*)
- ACID property in multi-user *transactions*

Terminology in RDB

Relational Schema

STUDENT (**stud#**, sname, address, deptno)
 DEPARTMENT (**deptno**, dname, chair, phone)



A Primary Key (PK) uniquely identifies each row.
 A Foreign Key (FK) is a set of attributes
 whose values are primary key values from another table.

Characteristics of Relational Model

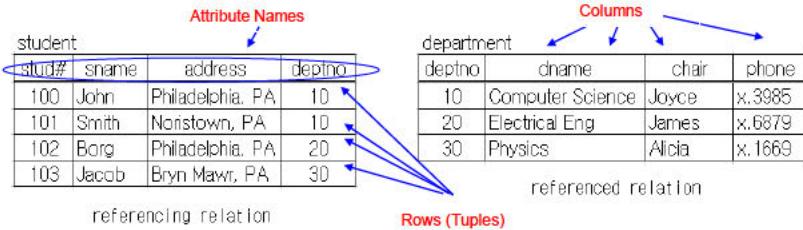


TABLE 3.1 Characteristics of a Relational Table

1	A table is perceived as a two-dimensional structure composed of rows and columns.
2	Each table row (tuple) represents a single entity occurrence within the entity set.
3	Each table column represents an attribute, and each column has a distinct name.
4	Each row/column intersection represents a single data value.
5	All values in a column must conform to the same data format.
6	Each column has a specific range of values known as the attribute domain .
7	The order of the rows and columns is immaterial to the DBMS.
8	Each table must have an attribute or a combination of attributes that uniquely identifies each row.

Example of a DDL in SQL Defining a Relational Schema

student			
stud#	sname	address	deptno
100	John	Philadelphia, PA	10
101	Smith	Norristown, PA	10
102	Borg	Philadelphia, PA	20
103	Jacob	Bryn Mawr, PA	30

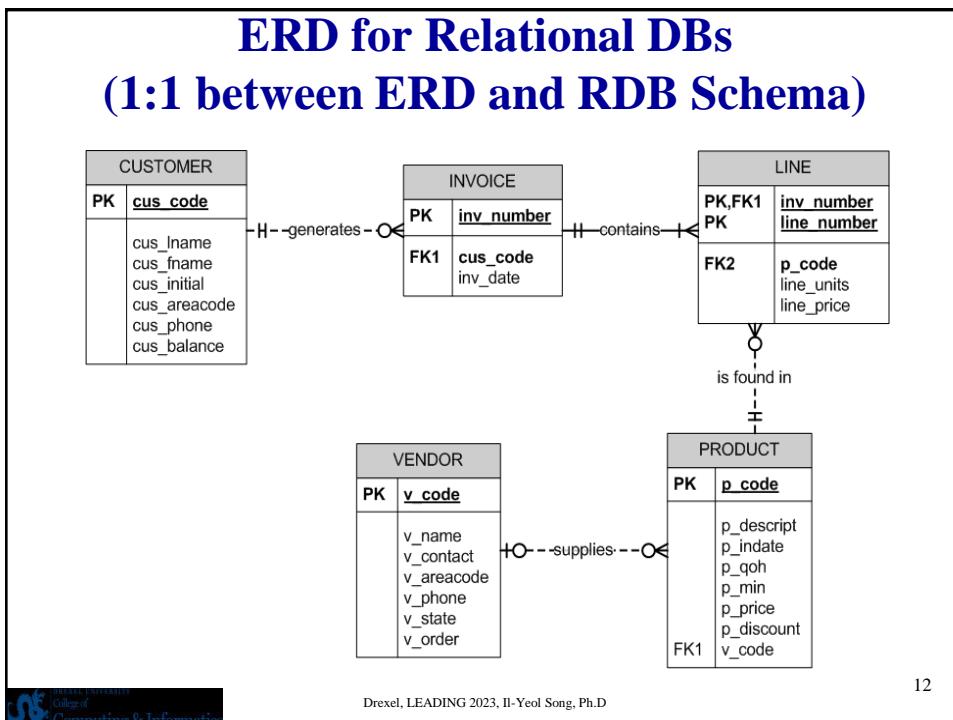
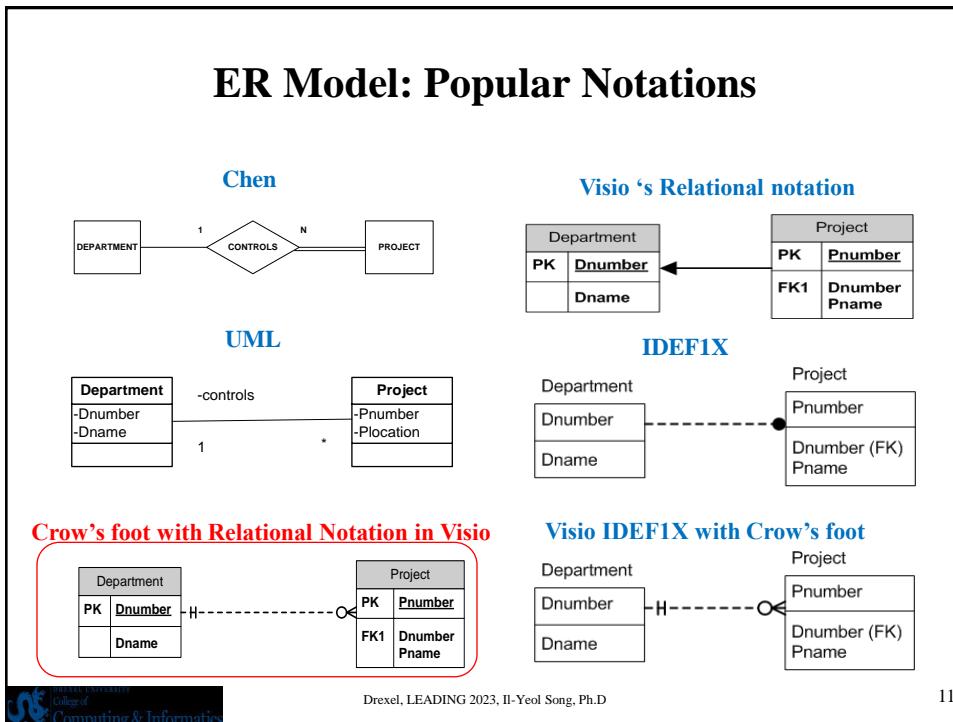
department			
deptno	dname	chair	phone
10	Computer Science	Joyce	x.3985
20	Electrical Eng	James	x.6879
30	Physics	Alicia	x.1669

referencing relation

```
CREATE TABLE Department (
    DeptNo      NUMBER(5)  PRIMARY KEY,
    DName       VARCHAR2(15) NOT NULL ,
    Chair        VARCHAR2 (20) ,
    Phone        CHAR (10) ,
);
```

```
CREATE TABLE Student (
    Stud#      NUMBER(8) ,
    SName       VARCHAR2(20) NOT NULL ,
    Address     VARCHAR2 (40) ,
    DeptNo      NUMBER (5)
    CONSTRAINT Stud_PK PRIMARY KEY (Stud#),
    CONSTRAINT Stud_FK FOREIGN KEY (DeptNo)
    REFERENCES Department (DeptNo)
);
```

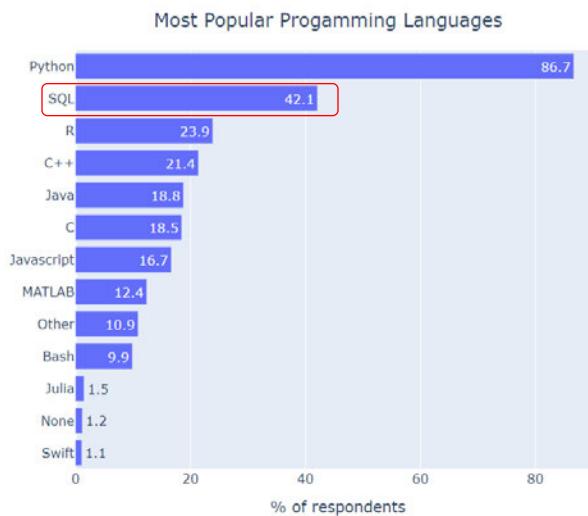
ER Model: Popular Notations



Why Relational Databases?

- Logical simplicity of the schema: **Tables**
- Easy, powerful, standard database language: **SQL**
- Transaction reliability: **ACID property** (Atomicity, Consistency, Isolation, Durability)
- Ad-hoc query processing
- Mature and reliable technologies
- Commercial investment over 45 years
- A large group of man-power and user bases

2020 Kaggle Data Science & Machine Learning Survey



Top Desired Skills in Data Science (2021)

11 Skills Data Scientists Need

1. Data Visualization
2. Python
- 3. SQL/NoSQL**
4. Social Media Mining
5. Fundamental Statistics
6. Natural Language Processing/Machine Learning
7. Microsoft Excel
8. High-Level Math
9. Teamwork
10. Communication
11. Business Savvy

Source: <https://bootcamp.berkeley.edu/blog/data-scientist-skills/>



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

15

15

Introduction to SQL

- Pronounced ‘see-quel’
- ANSI prescribes a standard SQL
- Developed by IBM in 1974, System R project
- First commercial implementation by Oracle (1976)
- ANSI/ISO Standard Versions:
 - SQL 86 (SQL 1)
 - SQL 89 (SQL 1 Level 2): RI integrity
 - SQL 92 (SQL 2)
 - SQL:1999 for OO features
 - SQL:2003 (XML features)
 - SQL: 2011 (Temporal data, windows functions)
 - SQL: 2016: JSON
 - SQL: 2019: Multi-dimensional Array (MDArray type and operations)
- Several dialects and extensions exist



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

16

16

SQL Topics

- **SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- **SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- **SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

17

17

Two major Components of SQL

- | Data Definition Language (DDL) | Data Manipulation Language(DML) |
|--|--|
| <ul style="list-style-type: none"> • CREATE TABLE • ALTER TABLE • DROP TABLE • CREATE/DROP VIEW • CREATE/DROP INDEX | <ul style="list-style-type: none"> • UPDATE TABLE • INSERT INTO • DELETE FROM • SELECT FROM |

DDL manages metadata
DML manages data



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

18

18

SQL Syntactic Rules

STUDENT

ID	Name	EnrollYear

- `SELECT enrollyear FROM student WHERE name = 'John';`
- SQL statement comprises
 - **Keywords**
 - Case-insensitive
 - Basic command set: fewer than 100 English words
 - Space-independent
 - **User-defined word (table names, attribute names)**
 - Case-insensitive
 - **Data**
 - **Case-sensitive** (In MS Access, case-insensitive)
 - String values: enclosed with a single quote
 - Ends with a semi-colon

An Overview of SQL Commands

DDL

--Changing attribute type

`ALTER TABLE project`

`MODIFY (budget
NUMBER(9,2));`

--Adding an attribute

`ALTER TABLE project ADD
(manager CHAR(10));`

-- Delete a column from a table

`ALTER TABLE project
DROP COLUMN
manager;`

--Removing a table from database

`DROP TABLE project;`

DML

--Inserting a row to a table

`INSERT INTO project VALUES (1234,
'Perfect Project', NULL, 'John');`

--Changing a value of attribute

`UPDATE project SET budget =
1.1*budget WHERE projno > 1000;`

--Deleting a row from a table

`DELETE FROM project WHERE
manager = 'John';`

Queries using SELECT Statement

SELECT	Specifies which columns are to appear in output
FROM	Specifies table(s) to be used
WHERE	Filters rows with conditions
GROUP BY	Forms groups of rows with the same column value.
HAVING	Filters groups subject to some condition.
ORDER BY	Specifies the order of the output.

- Only **SELECT** and **FROM** are mandatory.
- Order of the clauses cannot be changed.



An Overview of Simple Commands (I)

- (1) **SELECT * FROM Instructor;**
- (2) **SELECT DISTINCT fName FROM Instructor;**
- (3) **SELECT fName, lName, ssn FROM Instructor WHERE deptCode = 'math';**
- (4) **SELECT fName, lName FROM Instructor WHERE bonus > 1000;**
- (5) **SELECT * FROM Instructor WHERE (bonus >= 500 AND bonus <= 1000);**
- (6) **SELECT * FROM Instructor WHERE deptCode IS NULL;**
- (7) **SELECT * FROM Instructor WHERE deptCode IS NOT NULL;**
- (8) **SELECT * FROM Instructor WHERE bonus BETWEEN 500 AND 1000;**
- (9) **SELECT * FROM HR WHERE hireDate BETWEEN '01-MAY-2022' AND '31-MAY-2022';**
- (10) **SELECT * FROM Instructor WHERE bonus IN (100, 200, 300);**



An Overview of Simple Commands (II)

- (11) SELECT * FROM Instructor WHERE position = 'assistant';
- (12) SELECT * FROM R WHERE IName > 'S';
- (13) SELECT * FROM Instructor WHERE fName **LIKE** 'J%';
- (14) SELECT * FROM Instructor WHERE fName **LIKE** 'J_h_';
- (15) SELECT fName, lName FROM Instructor WHERE address **LIKE** '%Houston,TX%';
- (16) SELECT (salary+bonus) FROM Instructor;
- (17) SELECT (salary+bonus) **AS** total_income FROM Instructor;
- (18) SELECT ((A1/3.14)*2.54) **AS** A1_IN_CM FROM R;
- (19) SELECT **SYSDATE AS TODAY DUAL;**
- (20) **SELECT Fname, Lname, Bday,**
TRUNC(MONTHS_BETWEEN(SYSDATE, Bday)/12) AS "Actual Age"
FROM Person;



Join Operation

- JOIN is the relational operation that combines the data stored in two tables, creating a new result table
- JOIN is the most important operation in relational query processing
- **JOIN is performed between two tables that have a PK-FK relationship.**
 - If your FROM clause contains **2 tables**, you need **one join**.
 - If your FROM clause contains **3 tables**, you need **two joins**.
 - If your FROM clause contains **N tables**, you need **N-1 joins**.
- **JOINing between two tables that do not have a PK-FK relationship may NOT BE CORRECT.**



JOIN: Two Syntax

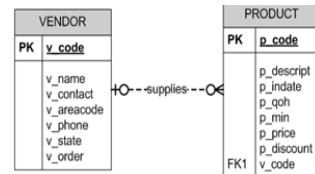
JOIN is usually done through ***PK-FK chains***

- For each product, find their vendor code and names

SELECT P.P_Code, V.V_Code, V.V_Name

FROM Product **P**, Vendor **V**

WHERE **P.V_Code = V.V_code;**



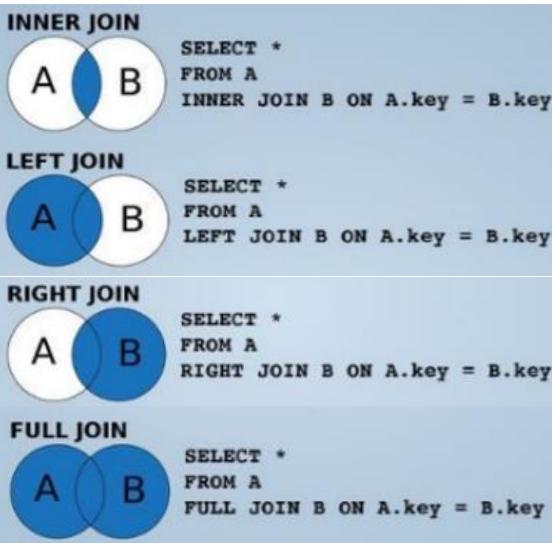
- New syntax**

SELECT P.P_Code, V.V_Code, V.V_Name

FROM Product **P INNER JOIN** Vendor **V**

ON P.V_Code = V.V_code;

Types of JOINS



NATURAL JOIN:

```
SELECT *
FROM table_A
NATURAL JOIN table_B;
```

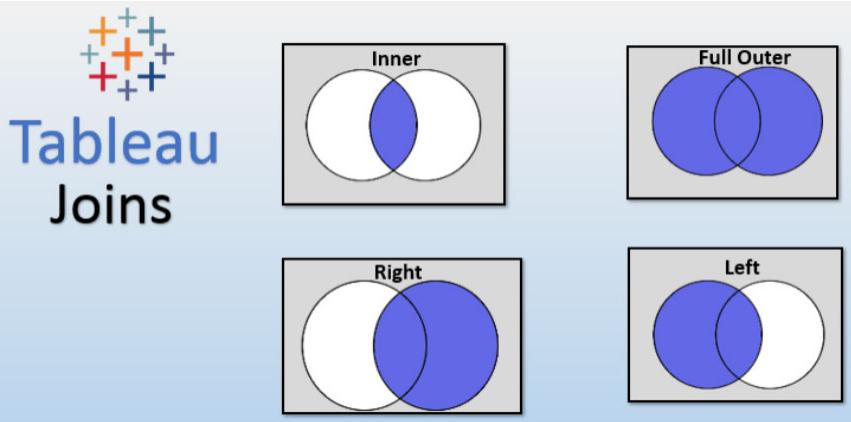
SELF JOIN:

```
SELECT *
FROM table_A X, table_A Y
WHERE X.A = Y.A;
```

CROSS JOIN:

```
SELECT *
FROM table_A
CROSS JOIN table_B;
```

Types of JOINS in Tableau



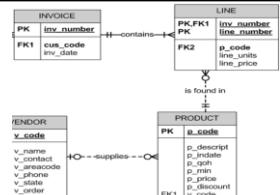
Drexel, LEADING 2023, Il-Yeol Song, Ph.D

27

27

(INNER) JOIN

- For each invoice, find the product description, #units, and price of items in the invoice.



```

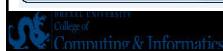
SQL> SELECT INVOICE.INV_NUMBER, PRODUCT.P_CODE, P_DESCRIP, LINE.UNITS, LINE.PRICE
2  FROM INVOICE JOIN LINE ON INVOICE.INV_NUMBER = LINE.INV_NUMBER
3  JOIN PRODUCT ON LINE.P_CODE = PRODUCT.P_CODE;

```

INV_NUMBER	P_CODE	P_DESCRIP	LINE_UNITS	LINE_PRICE
1001	13-Q2/P2	7.25-in. pwr. saw blade	1	14.99
1001	23109-HB	Claw hammer	1	9.95
1002	54778-2T	Rat-tail file, 1/8-in. fine	2	4.99
1003	2238/QPD	B&D cordless drill, 1/2-in.	1	38.95
1003	1546-QQ2	Hrd. cloth, 1/4-in., 2x50	1	39.95
1003	13-Q2/P2	7.25-in. pwr. saw blade	5	14.99
1004	54778-2T	Rat-tail file, 1/8-in. fine	3	4.99
1004	23109-HB	Claw hammer	2	9.95
1005	PVC23DRT	PVC pipe, 3.5-in., 8-ft	12	5.87
1006	SM-18277	1.25-in. metal screw, 25	3	6.99
1006	2232/QTY	B&D jigsaw, 12-in. blade	1	109.92
1006	23109-HB	Claw hammer	1	9.95
1006	89-WRE-Q	Hicut chain saw, 16 in.	1	256.99
1007	13-Q2/P2	7.25-in. pwr. saw blade	2	14.99
1007	54778-2T	Rat-tail file, 1/8-in. fine	1	4.99
1008	PVC23DRT	PVC pipe, 3.5-in., 8-ft	5	5.87
1008	WR3/TT3	Steel matting, 4'x8'x1/6", .5" mesh	3	119.95
1008	23109-HB	Claw hammer	1	9.95

18 rows selected.

SQL> -



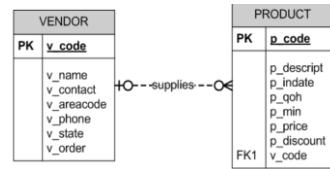
Drexel, LEADING 2023, Il-Yeol Song, Ph.D

28

28

LEFT (OUTER) JOIN

• **For ALL the Vendors**, find all the products they provide. Include all the vendors whether or not they provide any product.

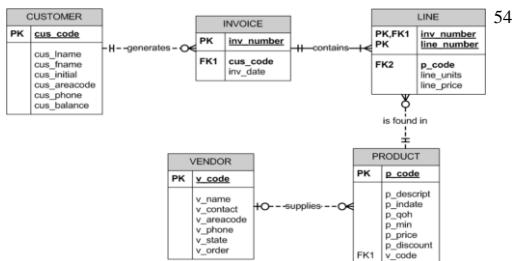


```
SQL> SELECT P_CODE, VENDOR.V_CODE, V_NAME
  2  FROM VENDOR LEFT JOIN PRODUCT ON VENDOR.V_CODE = PRODUCT.V_CODE;
```

P_CODE	V_CODE	V_NAME
11QER/31	25595	Rubicon Systems
13-02/P2	21344	Gomez Bros.
14-01/L3	21344	Gomez Bros.
1546-QQ2	23119	Randsets Ltd.
1558-QW1	23119	Randsets Ltd.
2232/QTY	24288	ORDVA, Inc.
2232/QWE	24288	ORDVA, Inc.
2238/QPD	25595	Rubicon Systems
23109-HB	21225	Bryson, Inc.
54778-2T	21344	Gomez Bros.
89-WRE-Q	24288	ORDVA, Inc.
SM-18277	21225	Bryson, Inc.
SW-23116	21231	D&E Supply
WR3/TT3	25595	Rubicon Systems
	22567	Dome Supply
	21226	SuperLoo, Inc.
	24004	Brackman Bros.
	25501	Damal Supplies
	25443	B&K, Inc.

19 rows selected.

Multi-table joins



List customer last name, invoice number, invoice date, and product description for all invoices for customer 10014.

```
SELECT      CUS_LNAME, I.INV_NUMBER, INV_DATE, P_DESCRPT
FROM        CUSTOMER C, INVOICE I, LINE L, PRODUCT P
WHERE       C.CUS_CODE = I.CUS_CODE
           AND LINV_NUMBER = L.INV_NUMBER
           AND L.P_CODE = P.P_CODE
           AND C.CUS_CODE = 10014
ORDER BY    INV_NUMBER;
```

Subqueries

- Subquery (or nested query): a complete **SELECT** statement embedded within another **SELECT** statement.

```

SELECT      select_list
FROM        table
WHERE       expr operator
            (SELECT      select_list
             FROM        table);
  
```

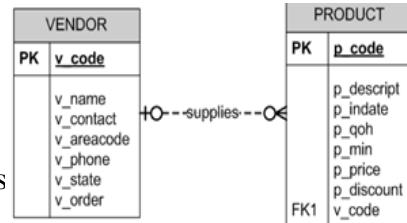
- The results of this inner **SELECT** statement are used in the outer **SELECT** statement to help determine the final result.
- A subquery produces a temporary table with results
- Can be used in **SELECT**, **FROM**, **WHERE**, and **HAVING** clauses of an outer **SELECT** statement and also in **INSERT**, **UPDATE**, **DELETE** statements

Subqueries

- Find all vendor names which supplies USB

```

SELECT v_name, v_phone
FROM Vendor
WHERE v_code IN (SELECT v_code
                  FROM Product
                  WHERE p_descript = 'USB');
  
```



- Can we use “=“ instead of “IN”?

```

SELECT v_name, v_phone
FROM Vendor
WHERE v_code = (SELECT v_code
                  FROM Product
                  WHERE p_descript = 'USB');
  
```

SQL Topics

- **SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- **SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- **SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

33

33

DATE Manipulation & Formats

- To insert a date using a default format:
`INSERT INTO Loan (Loan#, LoanDate) VALUES (12, '03-FEB-2019');`
- To insert a date using a non-default format:
`INSERT INTO Loan (Loan#, LoanDate)
VALUES (12, TO_DATE('06/02/2022', 'MM/DD/YYYY'));`
- Show today's date in the form of MM/DD/YYYY format
`SELECT SYSDATE, TO_CHAR(SYSDATE, 'MM/DD/YYYY') as
CurrDate FROM DUAL;`
- Show today's date including time
`SELECT SYSDATE, TO_CHAR(SYSDATE, 'yyyy-mm-dd hh24:mi:ss')
as CurrDateTime FROM DUAL;`



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

34

34

DATE Arithmetic

- Find products whose INDATE is more than 90 days old

```
SELECT      P_Code, P_Desc, P_Indate
FROM        Product
WHERE       P_Indate <= SYSDATE - 90;
```

- Find all the orders received in Feb 2023.

```
SELECT Order#, Odate
FROM ORDER
WHERE Odate BETWEEN '01-Feb-23' AND '28-Feb-23';
```



DATE Extraction

- **EXTRACT** function: returns numerical value of YEAR, MONTH, or DAY

Employee (E#, Ename, DoB)

-- Find the years of employees were born

```
SELECT      Ename, EXTRACT (YEAR FROM DoB)
FROM        Employee;
```

-- Month and day can also be extracted

```
SELECT Ename, EXTRACT (MONTH FROM DoB),
       EXTRACT (DAY FROM DoB)
FROM Employee;
```



STRING Manipulation

- Concatenating strings

SQL> `SELECT Lname||', '||Fname "FULL NAME" FROM Emp;`

will display

FULL NAME

Bond, James

- Handling Capitalization

SQL> `SELECT Lname, INITCAP(Lname), UPPER(Lname), LOWER(Lname) FROM Emp;`

- Capitalization: Show only the first letter in capital and the rest in lower

SQL> `SELECT INITCAP(LOWER(Lname)) FROM Emp;`

- Padding blanks to strings

RPAD: Pad to the right side of the column with left justification

LPAD: Pad to the left side of the column with right justification

SQL> `SELECT RPAD(Lname, 15, '.'), Age, LPAD(Lname, 15) FROM EMP;`

Clinton.....	53Clinton
--------------	----	--------------

Gore.....	52Gore
-----------	----	-----------

37



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

37

STRING Manipulation

/* TRIM function is used to remove all leading or trailing characters (or both) from a character string.*/

SQL> `SELECT TRIM(" LEADING Fellows ") FROM DUAL;`

/* Use of LTRIM(left trim) and RTRIM(right trim) functions*/

SQL> `SELECT Lname, LTRIM(Lname, 'SA'), RTRIM(Lname, 'S') FROM Emp;`

(LTRIM prints M from SAM, RTRIM prints WALES from WALES)

/* Remove . at the end and "THE from the front */

SQL> `SELECT LTRIM (RTRIM (Title,'.'), ' "THE ') FROM Magazine;`

Will convert

MY DARING."

"THE GOD FATHER."

into

MY DARING

GOD FATHER

38



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

38

STRING Manipulation

- **SUBSTR**

- Syntax : SUBSTR(char, m, n)
- Purpose : Returns a portion of the char, **beginning at character ‘m’ up to ‘n’ characters**. If ‘n’ is omitted, result is returned up to the end of char.
- Example : **SELECT SUBSTR(first_name,1,3)
FROM employee;**

Output :

```

SUB
---
JOH
KEV
JEA
LYN
LES
CYN

```



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

39

39

STRING Manipulation

```

/* SUBSTR(String, starting_position, #chars to be returned) */

/* Extract numeric part, add 1000, and concatenate */
SQL> SELECT Lname, E_ID,
      'S' || TO_CHAR( TO_NUMBER (SUBSTR (E_ID,2,3) ) + 1000)
      "NEW E_#"
   FROM Emp;

```

Will display

LNAME	E_ID	NEW E_#
Jones	E001	S1001
Wales	E002	S1002



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

40

40

NULL and NVL Function

- Changing the display of null values
SQL> SET NULL N/A
All null values are displayed as N/A.
SQL> SET NULL -- to reset the null to blank
- Using **NVL** function to change the null value to another value (Only in Oracle)
SQL> SELECT Fname, Lname, NVL(CreditCode, 999)
FROM Cust;

- If the value of CreditCode is not null, the value is displayed
- If the value of CreditCode is null, 999 is displayed.
- NVL(v1,v2): *v1* and *v2* must match by their data types.
- But *v1* can be any data type.

41

REPLACE() and TRANSLATE()

<https://www.databasestar.com/oracle-translate/>

TRANSLATE(string, from_string, to_string): From the *string*, replace each character of *from_string* into the corresponding character of the *to_string*, **char by char (case sensitive)**

REPLACE (string, from_string, to_string): From the *string*, replaces the **entire string** of *from_string* into *to_string*, rather than one at a time like TRANSLATE.

```
SELECT 'Complete IT Professional' as SAMPLE_TEXT,
REPLACE('Complete IT Professional', 'let', 'sip') as REPLACE_TEST,
TRANSLATE('Complete IT Professional', 'let', 'sip') as TRANSLATE_TEST
FROM dual;
```

Result:

SAMPLE_TEXT	REPLACE_TEST	TRANSLATE_TEST
Complete IT Professional	Compsipe IT Professional	Comsippi IT Profissionas

42

MERGE

<https://www.oracletutorial.com/oracle-basics/oracle-merge/>

- A command for copying/merging data from a source table to a target table, utilizing matching conditions

- Template:

```

MERGE INTO target_table
  USING source_table
  ON search_condition
  WHEN MATCHED THEN
    UPDATE SET col1 = value1, col2 = value2, ...
    WHERE <update_condition>
    [DELETE WHERE <delete_condition>]
  WHEN NOT MATCHED THEN
    INSERT (col1,col2,...)
    values(value1,value2,...)
    WHERE <insert_condition>;

```

43

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

43

MERGE

/* Update PRODUCT table with the data from NEW_PRODUCT
When their names matches, add the quantity; else insert a new row
with the new product info */

```

MERGE INTO product prod
  USING new_product new_prod
  ON (prod.product_name = new_prod.product_name)
  WHEN MATCHED THEN
    UPDATE SET prod.stock_qty = prod.stock_qty + new_prod.stock_qty
  WHEN NOT MATCHED THEN
    INSERT (product_name, stock_qty)
    VALUES (new_prod.product_name, new_prod.qty);

```

44

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

44

MERGE

- /* Some members got a new rank. Update them.
- Merge Member data to Member_Staging tables
 - When member_ID matches, but if Fname, Lname, or Rank are different, then use member table data. When member_ID does not match, insert */



• <https://www.oracletutorial.com/oracle-basics/oracle-merge/>

45

MERGE

<https://www.oracletutorial.com/oracle-basics/oracle-merge/>

```

MERGE INTO member_staging x
USING members y
ON (x.member_id = y.member_id)
WHEN MATCHED THEN
  UPDATE SET x.first_name = y.first_name,
             x.last_name = y.last_name,
             x.rank = y.rank
  WHERE x.first_name <> y.first_name OR
        x.last_name <> y.last_name OR
        x.rank <> y.rank
WHEN NOT MATCHED THEN
  INSERT(x.member_id, x.first_name, x.last_name, x.rank)
  VALUES(y.member_id, y.first_name, y.last_name, y.rank);

```

SQL Topics

- **SQL Basics**
 - Basic DDL and DML
 - Simple commands on a Single table
 - Joins (inner join, outer join, natural joins, multi-table joins)
 - Subquery
- **SQL for Data Engineering**
 - Date manipulation and Date Arithmetic
 - String manipulation
 - Merge
- **SQL for Analytics**
 - Aggregation, grouping, and ordering
 - OLAP functions
 - Ranking, partition by, buckets
 - Creating new derived columns



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

47

47

Analytic Functions in SQL

- SQL has many analytic functions:
 - Aggregation & Descriptive Statistics
 - OLAP (Online Analytical Functions)
 - Ranking
 - **RANK()** & **DENSE_RANK()**
 - Ranking within subgroups
 - **PARTITION BY**
 - Creating new derived columns
 - **WIDTH_BUCKET()**
 - **CASE...WHEN**
 - **COALESCE()**
 - **LAG()** and **LEAD()**
 - **CUME_DIST()**



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

48

48

Aggregate Functions

- Aggregation functions operate on a single column of a table and return a single value
- Five aggregate functions:
 - **COUNT**: returns the number of values in a specified column
 - **SUM**: returns the sum of the values in a specified column
 - **AVG**: returns the average of the values in a specified column
 - **MIN**: returns the smallest value in a specified column
 - **MAX**: returns the largest value in a specified column
- Where to use:
 - In the **SELECT** list
 - In the **HAVING** clause



Aggregation

- Example of five aggregation functions
SELECT MAX(bonus), MIN(bonus), AVG(bonus), COUNT(bonus), SUM(bonus)
FROM Instructor
WHERE position = ‘assistant’;
- How many **different** course titles are there?
SELECT COUNT(DISTINCT title) AS count
FROM Course;



Aggregation

- Note: What's wrong with the following?

```
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > AVG(bonus);
```

Note: use aggregate functions only in **SELECT** and **HAVING** clause

- **Why?**

- **Rewrite using a subquery**

```
SELECT instructorID, bonus
FROM Instructor
WHERE bonus > (SELECT AVG(bonus)
                FROM Instructor);
```



Descriptive Statistics in SQL

- **SELECT** cus_fname, cus_lname, cus_balance
AVG(cus_balance) mean,
MEDIAN (cus_balance) median,
STATS_MODE (cus_balance) mode,
STDDEV (cus_balance) “Standard Deviation”
FROM Product;

CUSTOMER	
PK	<u>cus_code</u>
	cus_lname
	cus_fname
	cus_initial
	cus_areacode
	cus_phone
	cus_balance

- **Computing median over partition**

```
SELECT cus_fname, cus_lname, cus_balance
MEDIAN (cus_balance) OVER (PARTITION BY cus_areacode) AS median
FROM Product;
```

- Computing SD over unique values

```
SELECT cus_fname, cus_lname, cus_balance
STDDEV (DISTINCT cus_balance) FROM Product;
```



Aggregation with GROUP BY

Find the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT deptCode, COUNT(instructorID) AS count, SUM(bonus) AS sum
FROM Instructor
GROUP BY deptCode
ORDER BY deptCode;
```

•Result:

deptCode	count	sum
acct	2	1100
math	2	300

Instructor
-instructorID{PK}
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus

- GROUP BY must include all non-aggregate function column names in the SELECT list.*



What's wrong with the following query?

Suppose we have the Instructor table as follows.

- Find the deptCode and deptName, and the total number of instructors in each department and the sum of their bonus, respectively

```
SELECT deptCode, deptName,
        COUNT(instructorID) AS count,
        SUM(bonus) AS sum
FROM Instructor
GROUP BY deptCode
ORDER BY deptCode;
```

Instructor
-instructorID{PK}
-fName
-lName
-ssn
-deptCode
-deptName
-position
-bonus



Aggregation with GROUP BY and HAVING

Oracle SQL*Plus

```
File Edit Search Options Help
SQL> SELECT V_CODE, COUNT(DISTINCT P_CODE), AVG(P_PRICE)
  2  FROM PRODUCT
  3  GROUP BY V_CODE;
```

V_CODE	COUNT(DISTINCT(P_CODE))	AVG(P_PRICE)
21225	2	8.47
21231	1	8.45
21344	3	12.49
23119	2	41.97
24288	3	155.593333
25595	3	89.63
	2	18.135

7 rows selected.

```
SQL> SELECT V_CODE, COUNT(DISTINCT P_CODE), AVG(P_PRICE)
  2  FROM PRODUCT
  3  GROUP BY V_CODE
  4  HAVING AVG(P_PRICE) < 10;
```

V_CODE	COUNT(DISTINCT(P_CODE))	AVG(P_PRICE)
21225	2	8.47
21231	1	8.45

PRODUCT	
PK	p_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

55

55

Aggregation with GROUP BY and HAVING

- Find the number of invoice per year and per month for the last 3 years

```
SELECT EXTRACT(year FROM inv_date) "Year",
       EXTRACT(month FROM inv_date) "Month",
       COUNT(inv_date) "No. of Invoices"
  FROM invoice
 GROUP BY EXTRACT(year FROM inv_date),
          EXTRACT(month FROM inv_date)
 HAVING EXTRACT(year FROM inv_date) IN (2017, 2016, 2015)
 ORDER BY "No. of Invoices" DESC;
```

INVOICE	
PK	inv_number
FK1	cus_code
	inv_date

56

Nested Subquery and Aggregate Functions

- Find the product that has the oldest date

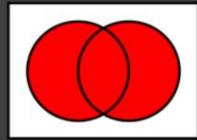
```
SELECT P_CODE, P_DESCRIP
FROM PRODUCT
WHERE P_INDATE = (
    SELECT MIN (P_INDATE)
FROM PRODUCT);
```

PRODUCT	
PK	p_code
	p_descrip
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code

This query is in the form of a nested query.

SET Operations

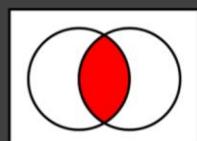
```
-- All customers and staff
SELECT first_name, last_name
FROM customer
UNION
SELECT first_name, last_name
FROM staff
```



UNION

R ∪ S

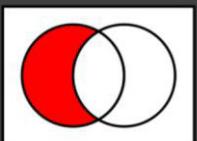
```
-- customers that are actors
SELECT first_name, last_name
FROM customer
INTERSECT
SELECT first_name, last_name
FROM actor
```



INTERSECT

R ∩ S

```
-- customers that are not staff
SELECT first_name, last_name
FROM customer
EXCEPT
SELECT first_name, last_name
FROM staff
```



MINUS

R - S

SQL OLAP Functions: ROLLUP

- Example: **GROUP BY ROLLUP (V_Code, P_Code)**
produces the union of
 - GROUP BY V_Code, P_Code
 - GROUP BY V_Code
 - Grand total

```
SQL> SELECT    U_CODE, P_CODE, SUM(TRADE_UNITS*TRADE_PRICE) AS TOTSALES
  2  FROM      DWTRADINGFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
  3  GROUP BY ROLLUP (U_CODE, P_CODE)
  4  ORDER BY U_CODE, P_CODE;
```

Oracle SQL*Plus

File Edit Search Options Help

```
SQL> SELECT    U_CODE, P_CODE, SUM(TRADE_UNITS*TRADE_PRICE) AS TOTSALES
  2  FROM      DWTRADINGFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWVENDOR
  3  GROUP BY ROLLUP (U_CODE, P_CODE)
  4  ORDER BY U_CODE, P_CODE;
```

U_CODE	P_CODE	TOTSALES
21225	23189-HB	99.5
21225	PUC23DRT	199.58
21225	SM-18277	41.94
21225		341.82
21344	13-Q2/P2	239.84
21344	54778-2T	59.88
21344		299.72
23119	1546-002	79.9
23119		79.9
24288	2232/QTY	219.84
24288	89-WRE-Q	513.98
24288		733.82
25595	2238/QPD	77.9
25595	WR3/TT3	719.7
25595		797.6
		2252.86

Subtotals by U_CODE

Grand total for all P_CODE values

16 rows selected.

Drexel, LEADING 2023, Il-Yeol Song, Ph.D.

SQL OLAP Functions: Cube

- The CUBE extension

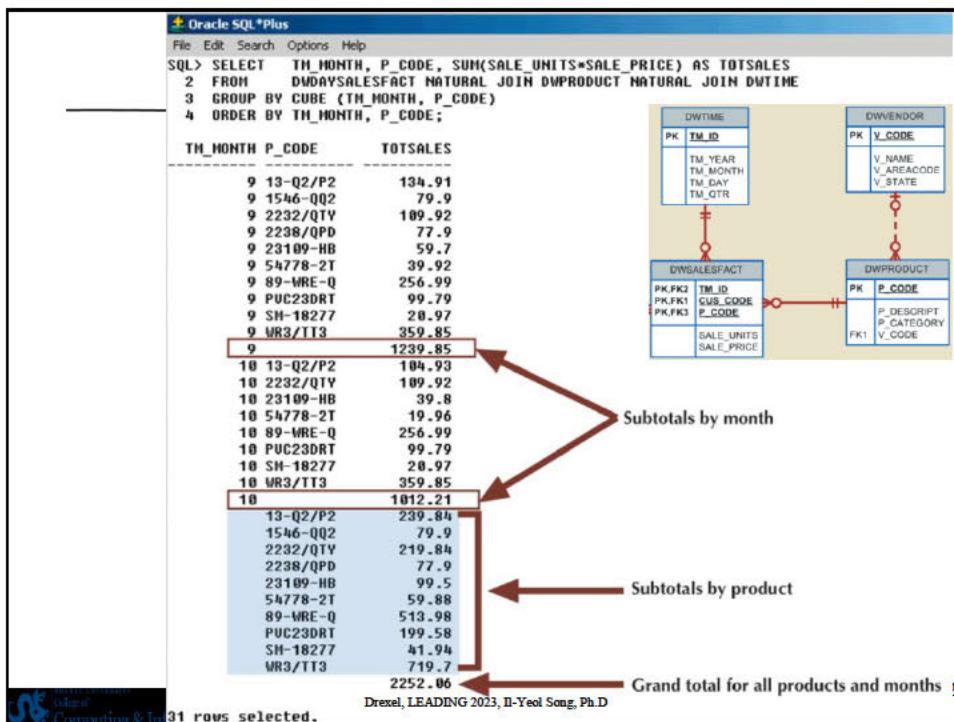
- Enables you to get a subtotal for each column listed in the expression, in addition to a grand total for the last column listed
- GROUP BY CUBE (TM_Month, P_CODE) is a union of:
 - GROUP BY (TM_Month, P_CODE)
 - GROUP BY (TM_Month)
 - GROUP BY (P_CODE)
 - Grand total

```
SQL> SELECT TM_MONTH, P_CODE, SUM( SALE_UNITS * SALE_PRICE ) AS TOTSALES
  2  FROM DW_DAYSALESFACT NATURAL JOIN DWPRODUCT NATURAL JOIN DWTIME
  3 GROUP BY CUBE ( TM_MONTH, P_CODE )
  4 ORDER BY TM_MONTH, P_CODE;
```

Drexel, LEADING 2023, Il-Yeol Song, Ph.D.

51

61



62

RANK and DENSE RANK Functions

- RANK() leave ranking gaps when there are multiple rows in the same rank (1, 2, 2, 4, 5...)
- DENSE_RANK() does **not** leave ranking gaps (1, 2, 2, 3, 4, 5)

```
SELECT p_code, p_descript, p_price,
       RANK() OVER
           (ORDER BY p_price NULLS LAST) AS Rank,
       DENSE_RANK()
           (ORDER BY p_price NULLS LAST) AS Dense_rank
  FROM Product
 ORDER BY p_price;
```

- DESC means descending order
- NULLS LAST means null values are smaller than non-null values
- You may use NULLS FIRST

PRODUCT	
PK	p_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code



RANK and DENSE RANK Functions

- Ranking applied to dates

```
SELECT p_code, p_descript, p_indate,
       RANK() OVER
           (ORDER BY p_indate NULLS LAST) AS Rank,
       DENSE_RANK()
           (ORDER BY p_indate NULLS LAST) AS Dense_rank
  FROM Product
 ORDER BY p_indate;
```

PRODUCT	
PK	p_code
	p_descript
	p_indate
	p_qoh
	p_min
	p_price
	p_discount
FK1	v_code



Top-N Queries

- In Oracle 12c and after:

```
SELECT      p_code, p_descript, p_price,
            RANK() OVER
                (ORDER BY p_price NULLS LAST) AS Rank,
FROM Product
ORDER BY p_price;
FETCH FIRST 10 ROWS ONLY;
```

- Another option by percentage

```
FETCH FIRST 20 PERCENT ROWS ONLY;
```

- MySQL uses LIMIT clause

```
LIMIT 10
```

65



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

65

PARTITION BY

- Partition the records by a column and then apply the function
- Ex: For each employee, show his/her dept# and the total# of employees working for the dept.

```
SELECT empno, deptno, COUNT(*)
       OVER (PARTITION BY deptno) DEPT_COUNT
  FROM emp;
```

emp_no	dept_no	DEPT_COUNT
1	10	3
2	10	3
3	10	3 <- three because there are three "dept_no = 10" records
4	20	2
5	20	2 <- two because there are two "dept_no = 20" records

<https://stackoverflow.com/questions/561836/oracle-partition-by-keyword>

66



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

66

RANK with PARTITION BY

- Create ranks within each partition

```
SELECT employee_id, first_name, department_id, salary  
RANK() OVER (PARTITION BY department_id  
          ORDER BY salary) rank  
FROM employee;
```

	EMPLOYEE_ID	FIRST_NAME	DEPARTMENT_ID	SALARY	RANK
1	119	Karen	30	2500	1
2	118	Guy	30	2600	2
3	117	Sigal	30	2800	3
4	116	Shelli	30	2900	4
5	115	Alexander	30	3100	5
6	114	Den	30	11000	6
7	144	Peter	50	2500	1
8	143	Randall	50	2600	2
9	142	Curtis	50	3100	3
10	141	Trenna	50	3500	4
11	107	Diana	60	4200	1
12	105	David	60	4800	2



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

67

67

Windows Function: WIDTH_BUCKET

SQL> SELECT lname, salary,

WIDTH_BUCKET(salary, 30000, 50000, 2) AS Bucket

FROM employee;

LNAME	SALARY	BUCKET
Smith	30000	1
Wong	40000	2
Zelaya	25000	0
Wallace	43000	2
Narayan	38000	1
English	25000	0
Jabbar	25000	0
Borg	55000	3



Drexel, LEADING 2023, Il-Yeol Song, Ph.D

68

68

CASE Statement

- Simple CASE Statements

```

SELECT      Fname, Lname,
(CASE      DNO
        WHEN 1      THEN 'Headquarters'
        WHEN 4      THEN 'Administration'
        WHEN 5      THEN 'Research'
        ELSE        'No department'
END) AS Department
FROM Employee;

```

Output:

Fname	Lname	Department
John	Smith	Research
Franklin	Wong	Research
Alica	Zelaya	Administration



CASE Statement

- Searched CASE Statements

```

SELECT      Fname, Lname, Salary
(CASE      Salary
        WHEN Salary <= 25000  THEN 1500
        WHEN Salary > 25000 AND Salary < 50000  THEN 1000
        WHEN Salary > 50000 AND Salary < 100000 THEN 500
        ELSE        0
END) AS "Bonus"
FROM Employee;

```

Output:

Fname	Lname	Salary	Bonus
John	Smith	30000	1000
Franklin	Wong	40000	1000
Alica	Zelaya	25000	1500



COALESCE Function

- A general form of NVL function-SQL 92 standard
- Provides a default value in case the column returns a NULL.

```

    S.storecode,
    coalesce(s.date, '2021-05-10') as date,
    coalesce(s.salesqty,0) as salesqty,
    coalesce(s.salesrevenue, 0)as salesrevenue
FROM Products P;
```

- NVL uses only 2 arguments, while COALESCE can have multiple arguments, but stops at the first NOT NULL value
- Example:

Gives a 10% discount to all products with a list price. If there is no list price, then the sale price is the minimum price. If there is no minimum price, then the sale price is "5":



COALESCE Function

- Gives a 10% discount to all products with a list price. If there is no list price, then the sale price is the minimum price. If there is no minimum price, then the sale price is "5":

```

SELECT product_id, list_price, min_price,
` COALESCE(0.9*list_price, min_price, 5) Sale
FROM product_information
WHERE supplier_id = 102050
ORDER BY product_id, list_price, min_price, Sale;
```

	PRODUCT_ID	LIST_PRICE	MIN_PRICE	Sale
	1769	48		43.2
	1770		73	73
	2378	305	247	274.5
Drexel, LEADING 2023, Il-Yeol Song, Ph.D	2382	850	731	765
	3355			5



LAG(), LEAD()

- Lead() allows you to grab data from *a next row* without having to do a self join and put it into the current row
- Lag() allows you to grab data from *a previous row* and put it into the current row
- Syntax:

```
LEAD(scalar_expression ,offset [,default])
OVER (
    [PARTITION BY partition_expression, ... ]
    ORDER BY sort_expression [ASC | DESC], ...
```



LAG(), LEAD()

- Display employee joinDate and EndingDate in one row

```
SELECT empcode, empname, joindate,
       LEAD (joinDate,1) OVER (ORDER BY joinDate ASC)
       as EndDate
FROM employee;
```

EMPCODE	EMPNAME	JOININGDATE
1	Rajendra	01-SEP-18
2	Manoj	01-OCT-18
3	Sonu	10-MAR-18
4	Kashish	25-OCT-18
5	Tim	01-DEC-18
6	Akshita	01-NOV-18

EMPCODE	EMPNAME	JOININGDATE	ENDDATE
3	Sonu	10-MAR-18	01-SEP-18
1	Rajendra	01-SEP-18	01-OCT-18
2	Manoj	01-OCT-18	25-OCT-18
4	Kashish	25-OCT-18	01-NOV-18
6	Akshita	01-NOV-18	01-DEC-18
5	Tim	01-DEC-18	- NVL



LAG(), LEAD()

- Given sales data below, display year, sales amount, and **increase from the previous year**

```
SELECT year, sales,
       Sale -LAG (sale, 1, NULL) OVER (ORDER BY year
ASC) AS increase-from-last-year
FROM sales_table;
```

year	sale	year	sale	increase-from-last-year
2017	100	2017	100	NULL
2018	125	2018	125	25
2019	140	2019	140	15
2020	175	2020	175	35



CUME_DIST()

- CUME_DIST() find the cumulative distribution of a value within a group of values, between 0 and 1.
 - Represents the number of rows that have values less than or equal to that row's value **divided by** the total number of rows

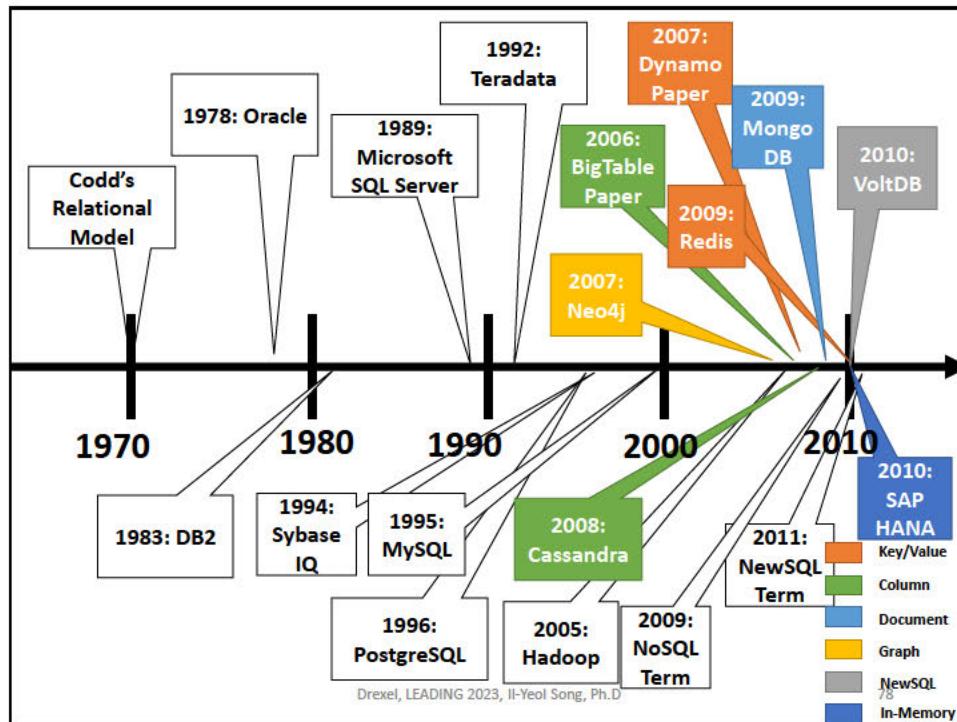
```
SELECT first_name, last_name, fees_paid,
CUME_DIST() OVER (ORDER BY fees_paid) AS cume_dist_test
FROM student;
```

FIRST_NAME	LAST_NAME	FEES_PAID	CUME_DIST_TEST
Julie	Armstrong	0	0.07142857143
Mark	Anderson	45	0.1428571429
Steven	Webber	80	0.2142857143
John	Smith	100	0.4285714286
Robert	Pickering	100	0.4285714286
Mary	Taylor	100	0.4285714286
Susan	Johnson	150	0.5714285714
Tanya	Hall	150	0.5714285714
Jarrad	Winston	300	0.6428571429
Tom	Capper	320	0.7142857143
Andrew	Cooper	400	0.7857142857
Mark	Holloway	410	0.8571428571
John	Rogers	700	0.9285714286
Michelle	Randall	(null)	1



Other Advanced SQL

- Other SQL functions and commands
- Views and Materialized views
- Indexes
- SQL with JSON
- SQL with Python
- PL/SQL
 - Stored functions
 - Stored Procedures
 - Triggers
 - Packages



Major Features of NoSQL Databases (1)



- Manage ***non-relational data***: Semi and non-structured data
- Stores ***aggregated objects***:
 - Stores all the related data together (denormalized objects)
 - No complex relationships
 - Could cause redundant data and update anomaly
 - Write-once-read-many applications
- ***Schema-on-read, mostly***
 - Some column stores need to define the schema in advance, but columns can be added dynamically
- Built on ***the scale-out architecture***
 - Parallel processing, Scalability, fault tolerance, High availability (HA)
 - Support ***automatic sharding***
 - Data are automatically distributed to nodes based on some fields

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

79

79

Major Features of NoSQL Databases (2)



- Soften the ACID property:
 - **BASE**: *Basically Available, Soft state, Eventual consistency*
 - **NoSQLs focus on scalability and availability**
 - The data (and its replicas) will become consistent at some point in time after each transaction
 - Supports high throughput and low latency
- Open-sourced
- Four types
 - Key-value stores, Wide Column stores, Document stores
 - Graph databases
- Designed for ***real-time OLTP system for non-uniform, big data.***

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

80

80

Rankings of DB Systems by Popularity

- NoSQL is moving up.
 - DB-Engines Ranking by popularity (<https://db-engines.com/en/ranking> May, 2022): From Websites, Google trends, social networks, job Ads, etc.

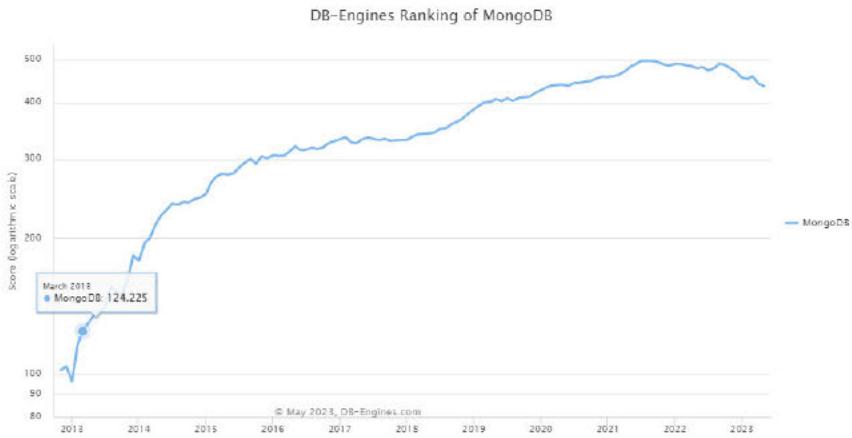
Rank	DBMS	Database Model		
May 2023	Apr 2023	May 2022	DBMS	Database Model
1.	1. Oracle	Relational, Multi-model		
2.	2. MySQL	Relational, Multi-model		
3.	3. Microsoft SQL Server	Relational, Multi-model		
4.	4. PostgreSQL	Relational, Multi-model		
5.	5. MongoDB	Document, Multi-model		
6.	6. Redis	Key-value, Multi-model		
7.	7. IBM DB2	Relational, Multi-model		
8.	8. Elasticsearch	Search engine, Multi-model		
9.	9. SQLite	Relational		
10.	10. Microsoft Access	Relational		
11.	11. Snowflake	Relational		
12.	12. Cassandra	Wide column		
13.	13. MariaDB	Relational, Multi-model		
14.	14. Splunk	Search engine		
15.	15. Amazon DynamoDB	Multi-model		
16.	16. Microsoft Azure SQL Database	Relational, Multi-model		
17.	17. Hive	Relational		
18.	18. Databricks	Multi-model		
19.	19. Teradata	Relational, Multi-model		
20.	20. Google BigQuery	Relational		
21.	21. FileMaker	Relational		
22.	22. Neo4j	Graph		
23.	23. SAP HANA	Relational, Multi-model		

81
Drexel LEADING 2023, Il-Yeol Song, Ph.D

81

MongoDB Popularity score

https://db-engines.com/en/ranking_trend/system/MongoDB
(May 2023)



82

82

MongoDB Atlas

- A cloud-based MongoDB cluster (A MongoDB as a Service)
- <https://www.mongodb.com/atlas/database>
- [Deploy a Free Cluster — MongoDB Atlas](#)
- 512MB free storage

Cloud Provider & Region

AWS, N. Virginia (us-east-1)

 Multi-Cloud, Multi-Region & Workload Isolation (M10+ clusters)
Distribute data across clouds    or regions for improved availability and local read performance, or introduce replicas for workload isolation. [Learn more](#)

OFF

Drexel University College of Computing & Informatics

Drexel, LEADING 2023, Il-Yeol Song, Ph.D

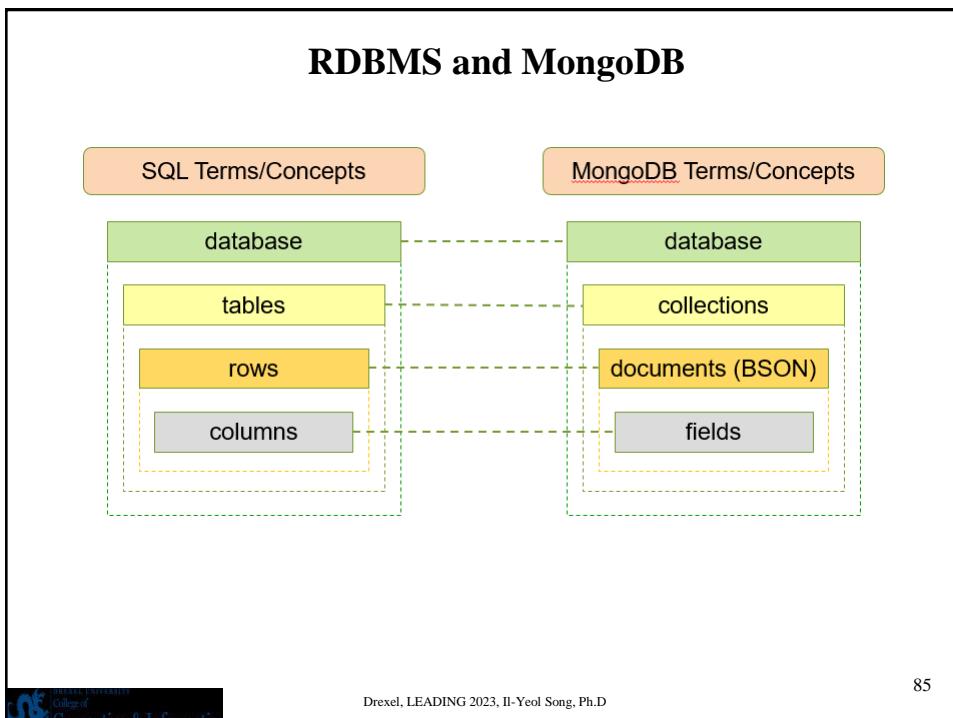
83

MongoDB Introduction



- MongoDB is from “**Humongous**”
- A **document-based** NoSQL database by MongoDB, Inc.
 - Documents are stored as **JSON** files; this makes it easier to read and manipulate using different programming languages.
 - Stored in **BSON** format for efficient storage, indexing, and querying
 - Optimized for *queries* of **semi-structured data**
 - Used when scalability, flexible schema design, and real-time data access (as in IOT, real-time analytics) are important,
 - Horizontal scaling (fault-tolerance, replication, high availability)
 - Schemaless
 - Released in 2009.

RDBMS and MongoDB



JSON

- JavaScript Object Notation (JSON) – data interchange format used to represent data as a logical object.
- Every JSON document requires an object ID.
- **Object** is enclosed by a pair of **curly brackets** {K:V}
- **Array** is enclosed by a pair of **square brackets** []
- **Value could be a “string”, an object, an array of values, a set of objects.**

```
{
  _id: 101,
  title: "Database Systems",
  author: ["Coronel", "Morris"],
  publisher: {
    name: "Cengage",
    street: "500 Topbooks Avenue",
    city: "Boston",
    state: "MA"
  }
}
```

- Example:

CRUD Example

```
> db.friends.insertOne({
    first: "John",
    last : "Doe",
    cell: "215-123-4567",
    city: "Philly"
})
```

```
> db.friends.find (
    {city : "Philly"},
    {name:1, cell: 1, _id:0 }
)
```

```
> db.friends.updateOne(
    {first:"John", last: "Doe"},
    {$set: {
        cell: "215-456-7890",
        type: "personal"
    }
})
```

```
>db.friends.deleteMany(
    { city : "Boston" }
)
```

CRUD: Creating/Inserting documents

- Example:- `db.<<collectionName>>.insertOne({document})`

```
db.users.insertOne( ← collection
{
    name: "sue", ← field: value
    age: 26, ← field: value
    status: "pending" ← field: value
}
) } document
```

```
>>> db.emp.insert({id: 1, name: "Song"})
WriteResult({ "nInserted" : 1 })
>>> |
```

CRUD: Inserting Multiple documents

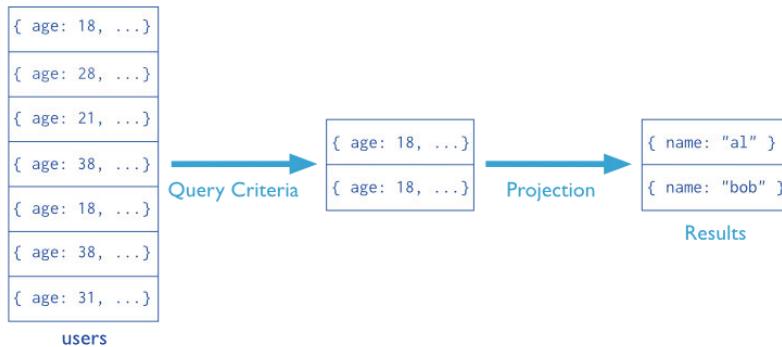
- Use the method – **insertMany([**{d1}, {d2}.., {dn}**])**
- Example:-

```
db.inventory.insertMany([
  { item: "journal", qty: 25, tags: ["blank", "red"], size: { h: 14, w: 21, uom: "cm" } },
  { item: "mat", qty: 85, tags: ["gray"], size: { h: 27.9, w: 35.5, uom: "cm" } },
  { item: "mousepad", qty: 25, tags: ["gel", "blue"], size: { h: 19, w: 22.85, uom: "cm" } }
])
```

89

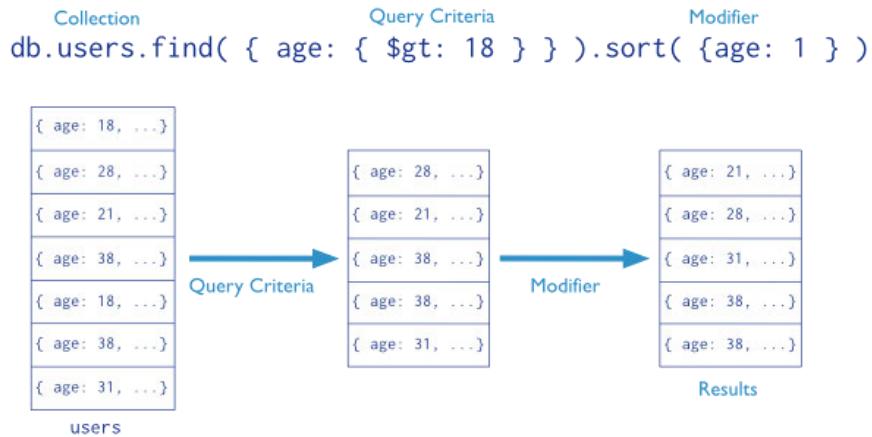
CRUD: Find() with Query and Projections

Collection	Query Criteria	Projection
<code>db.users.find({ age: 18 }, { name: 1, _id: 0 })</code>		



90

CRUD: find() with Query Criteria



CRUD: UPDATE

- Update ***all the documents*** whose age is greater than 25

```
UPDATE people
SET status = "C"
WHERE age > 25
```

```
db.people.updateMany(
  { age: { $gt: 25 } },
  { $set: { status: "C" } }
)
```

- Update ***the first document*** whose age is greater than 25

```
db.people.updateOne(
  { age: { $gt: 25 } },
  { $set: { status: "C" } }
)
```

CRUD: DELETE

```
DELETE FROM users
WHERE status = "D"
```

- Remove ***the first*** document from the users collection where the status field equals “D”:

```
db.users.deleteOne( { status: "D" } )
```

- Remove ***all the documents*** from the users collection where the status field equals “D”:

```
db.users.deleteMany({ status : "D" })
```

- Remove ***all the documents*** from the users collection:

```
db.users.deleteMany({ })
```



MongoDB: Other Topics

- Pattern Matching
- Array manipulation
- Aggregation framework
- Indexing
- Document design
- PyMongo
- Architecture
- Replication
- Sharding



Good Resources

- **SQL Tutorial - Full Database Course for Beginners,**
freeCodeCamp.org,
<https://www.youtube.com/watch?v=HXV3zeQKqGY>
- **SQL Tutorial,** <https://www.w3schools.com/sql/>
- **Mongo DB Basics,** [Jyoti Reddy, Data Engineer at Krones](#)
<https://www.kdnuggets.com/2019/06/mongo-db-basics.html>
- **MongoDB Tutorial for Beginners: Learn Basics in 7 Days,**
[David Taylor](#) April 23, 2022.
<https://www.guru99.com/mongodb-tutorials.html>



Summary

- Relational databases are used to manage **record-oriented structured data** for the applications in which transaction reliability (ACID) is important.
 - SQL is the standard DB language for RDBs
- NoSQL databases are used for **managing non-relational data** in the form of key-value pair, wide columns, unstructured documents, or graph-structure.
- NoSQL databases offer **flexibility, scalability, and performance** advantages in certain big data use cases where traditional relational databases may not be the best fit
- MongoDB is the most widely used **document databases** that manages unstructured or semi-structured data in the form of JSON format
 - Schemaless and No SQL support yet (mostly)



Question?

